



# From Thousands of Hours to a Couple of Minutes: Towards Automating Exploit Generation for Arbitrary Types of Kernel Vulnerabilities

# Who are We?



8/9/18

- Wei Wu [@wu\\_xiao\\_wei](https://twitter.com/wu_xiao_wei)
  - Visiting scholar at JD.com
    - Conducting research on software security in Enterprise Settings
  - Visiting Scholar at Penn State University
    - Vulnerability analysis
    - Memory forensics
    - Malware dissection
    - Reverse engineering
    - Symbolic execution
    - Static analysis
  - Final year PhD candidate at UCAS
    - Knowledge-driven vulnerability analysis
  - Co-founder of CTF team Never Stop Exploiting.(2015)
    - ctftime 2017 ranking 4<sup>th</sup> team in China
  - I am on market.

## NSA Codebreaker Challenge

University

Carnegie Mellon University

Lafayette College

University of Hawaii

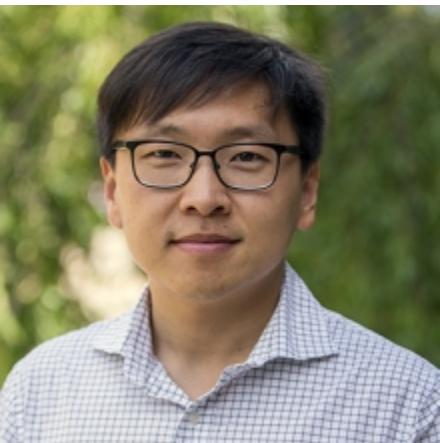
Pennsylvania State University

Georgia Institute of Technology

## China

Position	Country position	Name	Points	Events
21	1	eee	365.576	16
24	2	A*0*E	332.972	6
27	3	0ops	278.460	18
30	4	Never Stop Exploiting	247.073	9
34	5	Azure Assassin Alliance	235.876	25

# Xinyu Xing



- Visiting scholar at JD.com
  - Conducting research on software and hardware security in Enterprise Settings
- Assistant Professor at Penn State University
  - Advising PhD students and conducting many research projects on
    - Vulnerability identification
    - Vulnerability analysis
    - Exploit development facilitation
    - Memory forensics
    - Deep learning for software security
    - Binary analysis
    - ...

# • Jimmy Su



## Head of JD security research center

- Vulnerability identification and exploitation in Enterprise Settings
- Red Team
- JD IoT device security assessments
- Risk control
- Data security
- Container security

# What are We Talking about?



#BHUSA

- Discuss the challenge of exploit development
- Introduce an automated approach to facilitate exploit development
- Demonstrate how the new technique facilitate mitigation circumvention

# Background



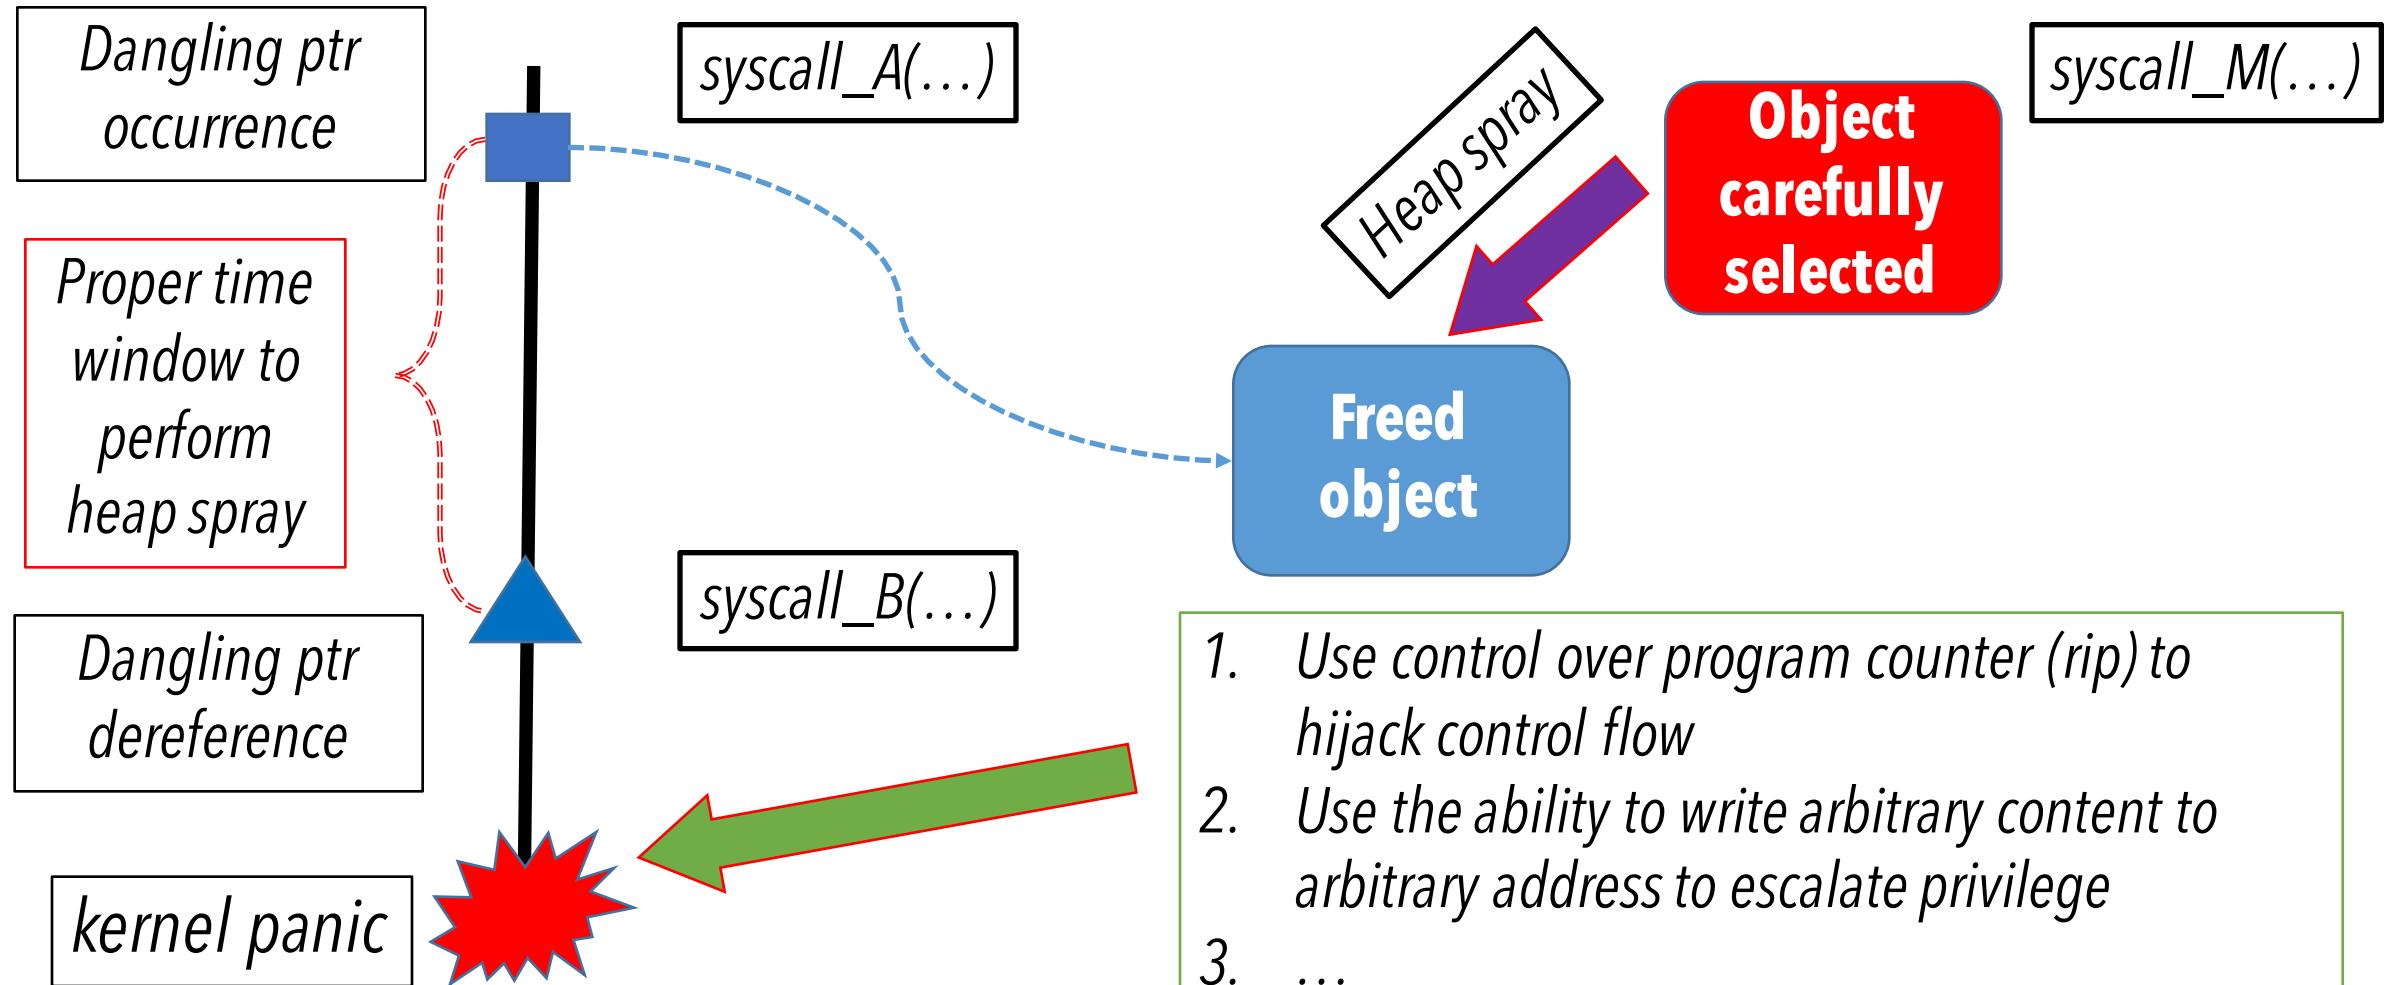
- All software contain bugs, and # of bugs grows with the increase of software complexity
  - E.g., Syzkaller/Syzbot reports 800+ Linux kernel bugs in 8 months
- Due to the lack of manpower, it is very rare that a software development team could patch all the bugs timely
  - E.g., A Linux kernel bug could be patched in a single day or more than 8 months; on average, it takes 42 days to fix one kernel bug
- The best strategy for software development team is to prioritize their remediation efforts for bug fix
  - E.g. based on its influence upon usability
  - E.g., based on its influence upon software security
  - E.g., based on the types of the bugs
  - ....

# Background (cont.)



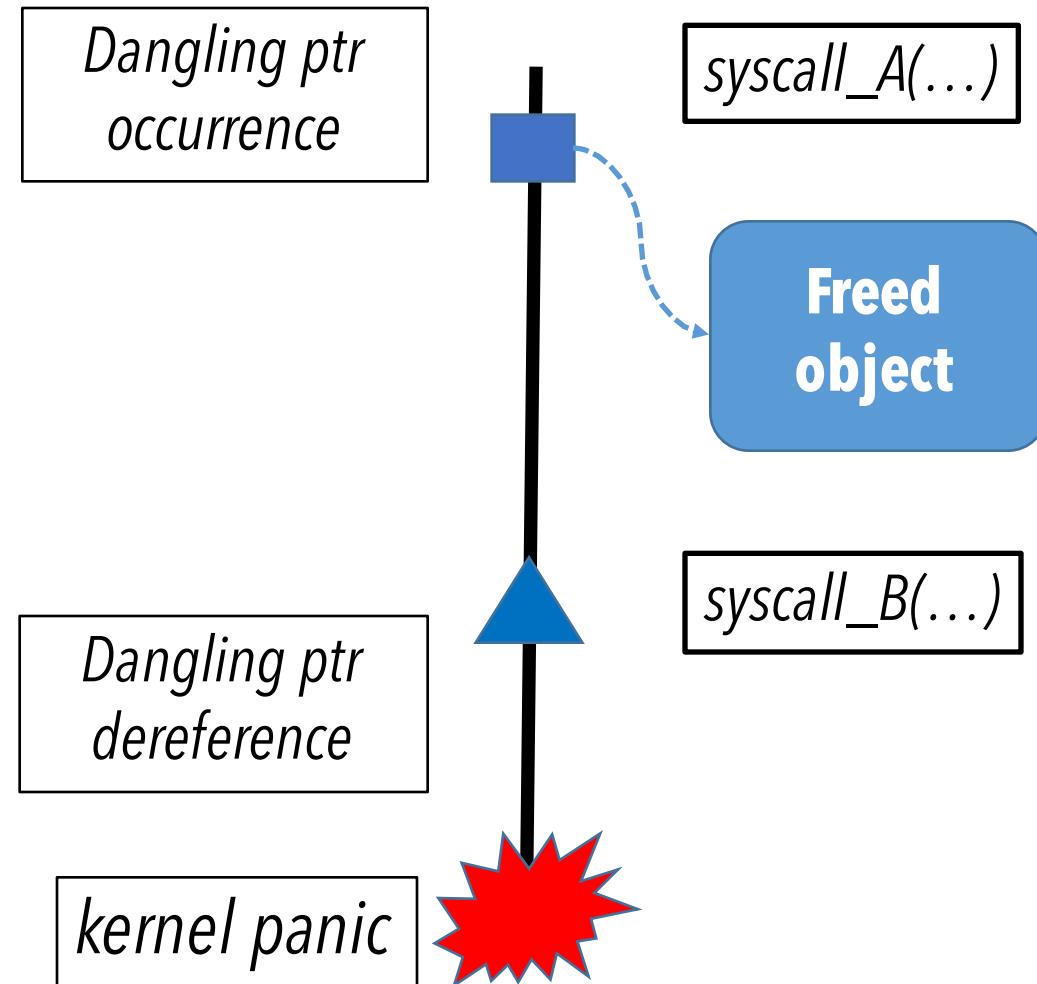
- Most common strategy is to fix a bug based on its exploitability
- In the blue team, you need to “attack” with exploitable vulnerabilities
- To determine the exploitability of a bug, analysts generally have to write a working exploit, which needs
  - 1) Significant manual efforts
  - 2) Sufficient security expertise
  - 3) Extensive experience in target software

# Crafting an Exploit for Kernel Use-After-Free



# Challenge 1: Needs Intensive Manual Efforts

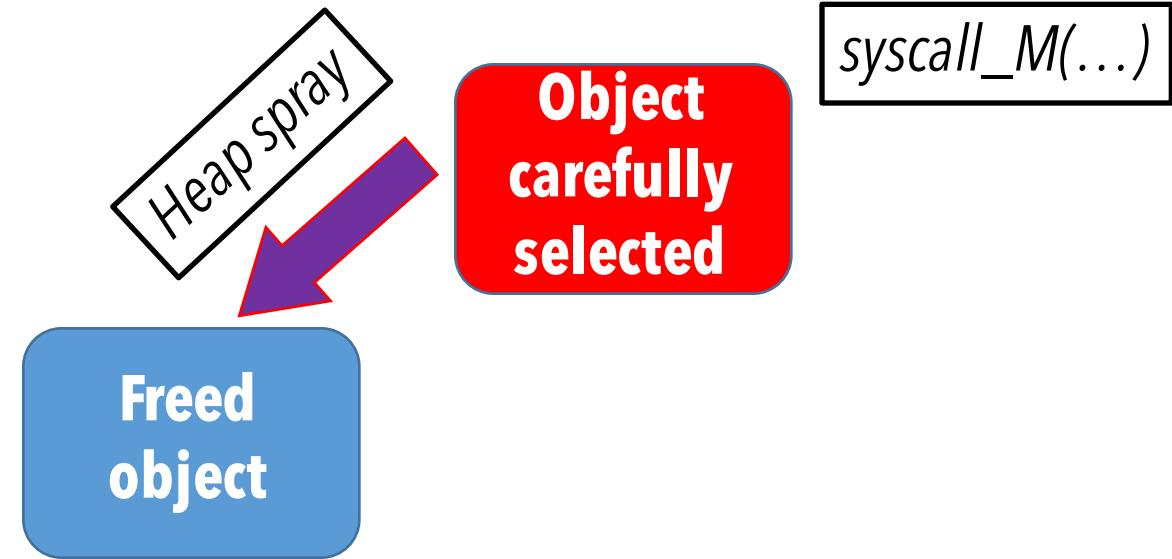
- Analyze the kernel panic
- Manually track down
  1. The site of dangling pointer occurrence and the corresponding system call
  2. The site of dangling pointer dereference and the corresponding system call



# Challenge 2: Needs Extensive Expertise in Kernel

#BHUSA

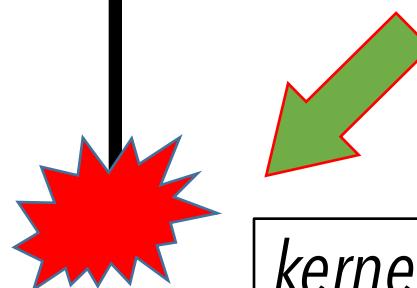
- Identify all the candidate objects that can be sprayed to the region of the freed object
- Pinpoint the proper system calls that allow an analyst to perform heap spray
- Figure out the proper arguments and context for the system call to allocate the candidate objects



# Challenge 3: Needs Security Expertise

- Find proper approaches to accomplish arbitrary code execution or privilege escalation or memory leakage
  - E.g., chaining ROP
  - E.g., crafting shellcode
  - E.g., fixing critical data
  - ...

1. *Use control over program counter (rip) to perform arbitrary code execution*
2. *Use the ability to write arbitrary content to arbitrary address to escalate privilege*
3. ...



*kernel panic*

- Approaches for Challenge 1
  - Nothing I am aware of, but simply extending KASAN could potentially solve this problem
- Approaches for Challenge 2
  - [Blackhat07][Blackhat15][USENIX-SEC18]
- Approaches for Challenge 3
  - [NDSS'11] [S&P16], [S&P17]

[NDSS11] Avgerinos et al., AEG: Automatic Exploit Generation.

[Blackhat 15] Xu et al., Ah! Universal android rooting is back.

[S&P16] Shoshtaishvili et al., Sok:(state of)the art of war: Offensive techniques in binary analysis.

[USENIX-SEC18] Heelan et al., Automatic Heap Layout Manipulation for Exploitation.

[S&P17] Bao et al., Your Exploit is Mine: Automatic Shellcode Transplant for Remote Exploits.

[Blackhat07] Sotirov, Heap Feng Shui in JavaScript

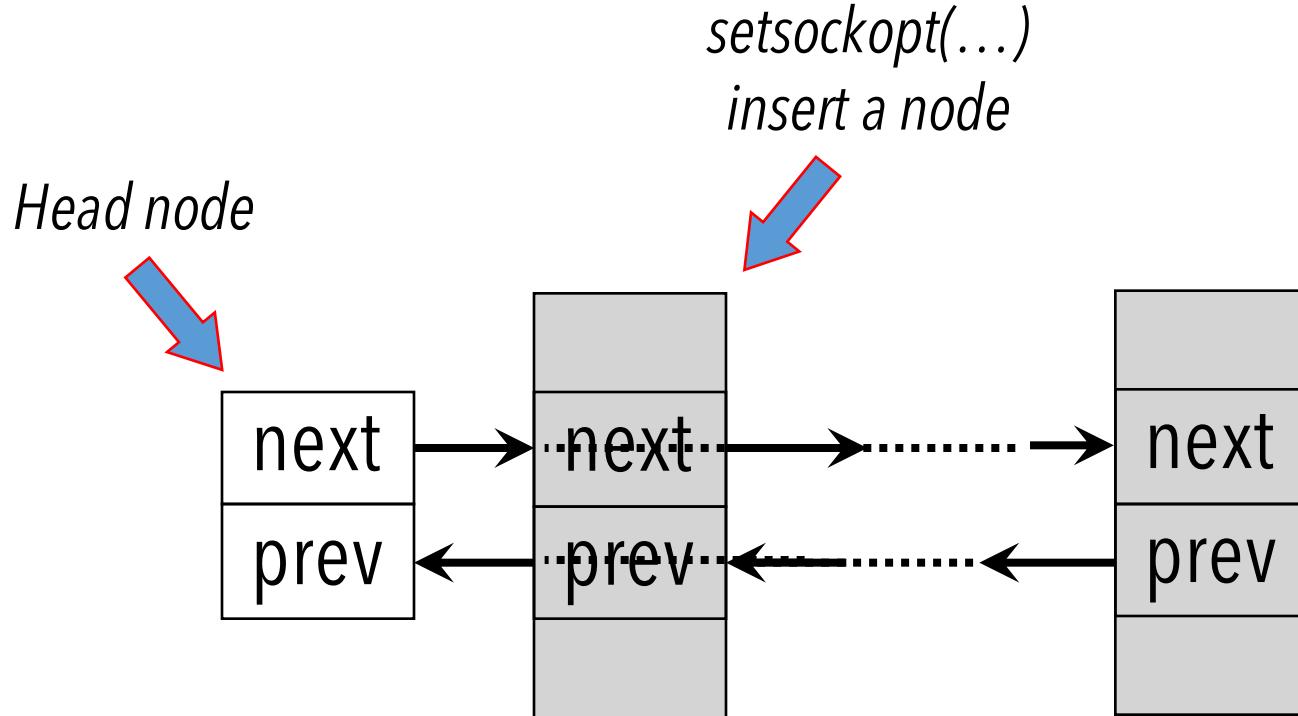


# Roadmap



- Unsolved challenges in exploitation facilitation
- Our techniques -- FUZE
- Evaluation with real-world Linux kernel vulnerabilities
- Conclusion

# A Real-World Example (CVE-2017-15649)



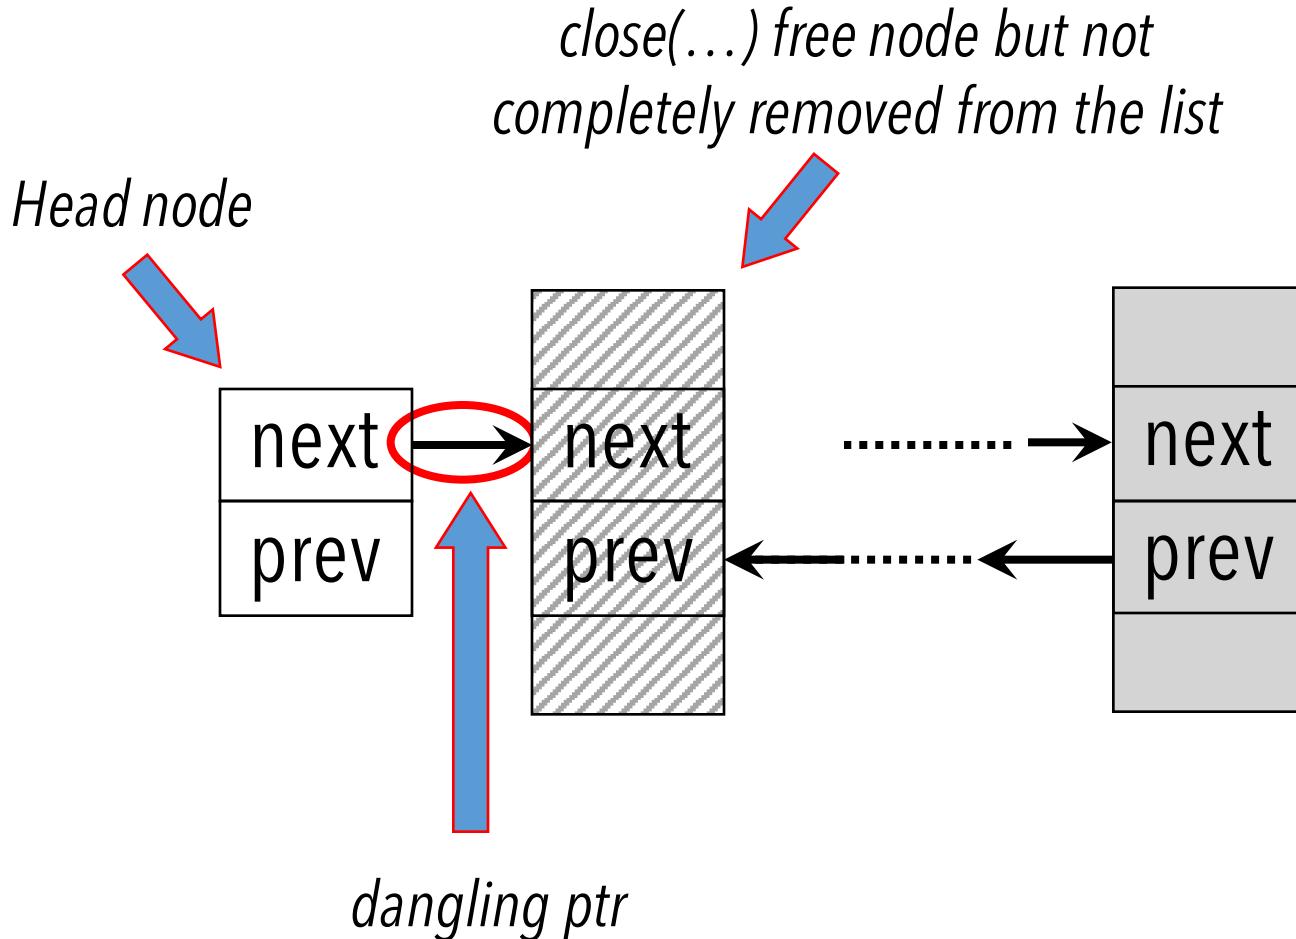
```

1 void *task1(void *unused) {
2 ...
3     int err = setsockopt(fd, 0x107, 18,
4     ↳ ... , ... );
5 }
6 void *task2(void *unused) {
7     int err = bind(fd, &addr, ... );
8 }
9
10 void loop_race() {
11 ...
12     while(1) {
13         fd = socket(AF_PACKET, SOCK_RAW,
14         ↳ htons(ETH_P_ALL));
15 ...
16     //create two racing threads
17     pthread_create(&thread1, NULL,
18     ↳ task1, NULL);
19     pthread_create(&thread2, NULL,
20     ↳ task2, NULL);
21
22     pthread_join(thread1, NULL);
23     pthread_join(thread2, NULL);
24 }
25 }
```

Annotations:

- Line 3: `setsockopt(fd, 0x107, 18, ... , ... );` - Circled in green.
- Line 13: `fd = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_ALL));` - Circled in red.
- Line 16: `pthread_create(&thread1, NULL, task1, NULL);` - Circled in green.
- Line 22: `close(fd);` - Circled in red.

# A Real-World Example (CVE 2017-15649)

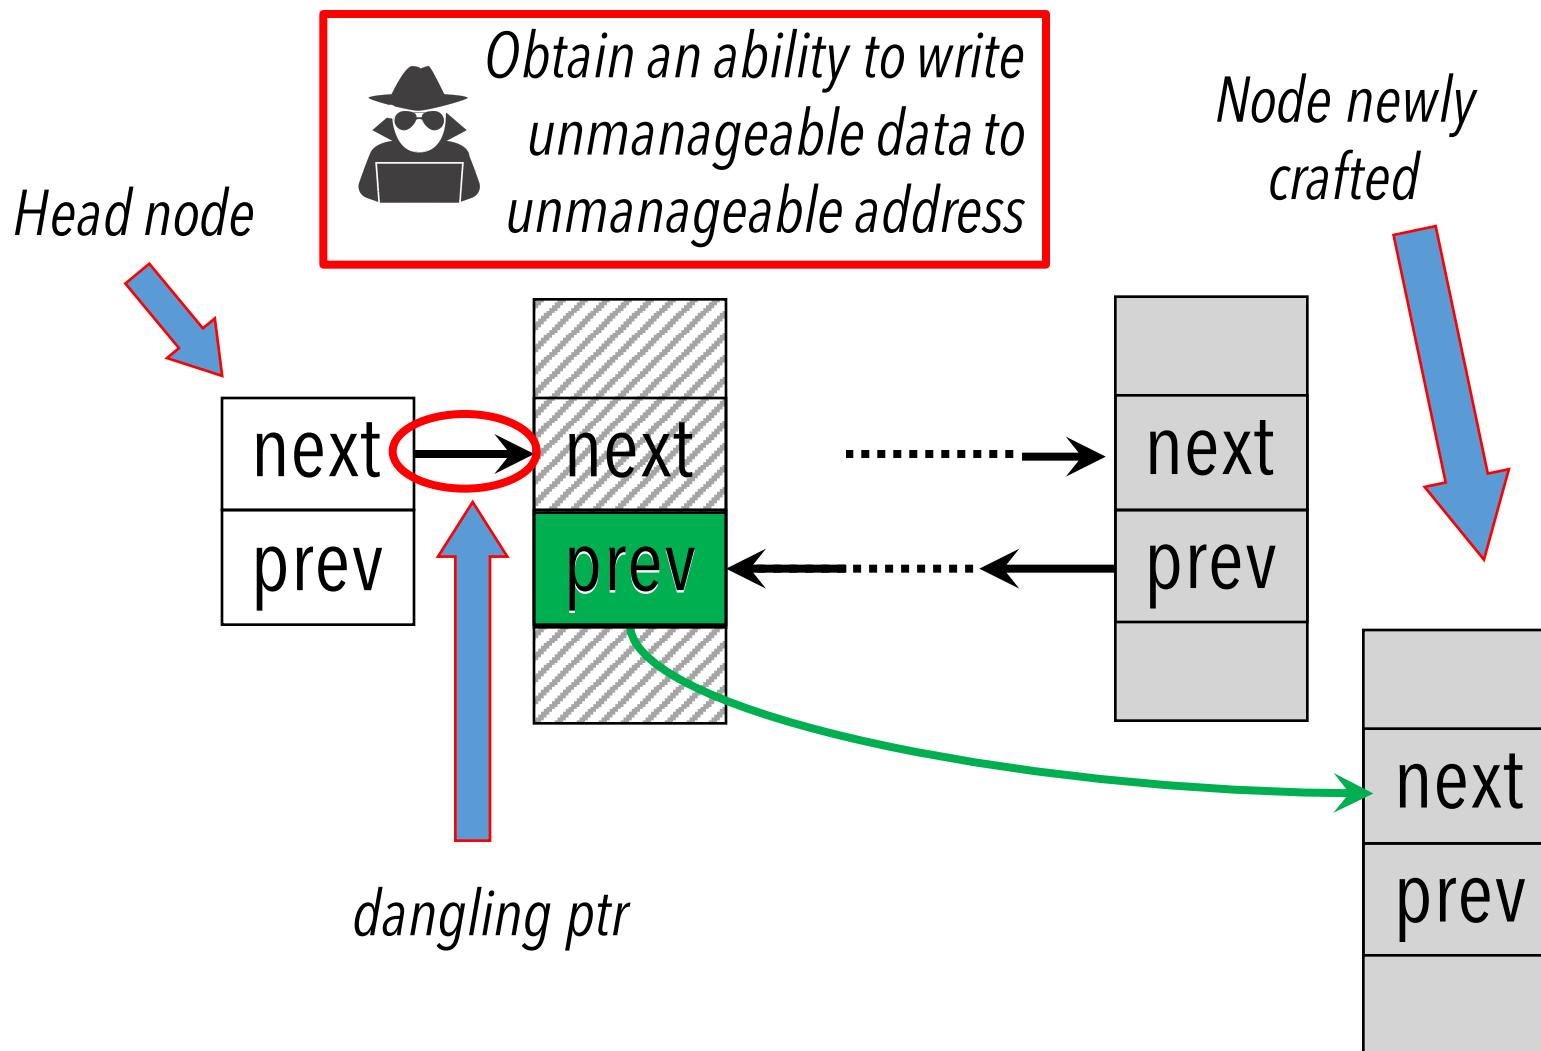


```

1 void *task1(void *unused) {
2 ...
3     int err = setsockopt(fd, 0x107, 18,
4     ↪ ..., ...);
5 }
6 void *task2(void *unused) {
7     int err = bind(fd, &addr, ...);
8 }
9 void loop_race() {
10 ...
11 while(1) {
12     fd = socket(AF_PACKET, SOCK_RAW,
13     ↪ htons(ETH_P_ALL));
14 ...
15 //create two racing threads
16 pthread_create(&thread1, NULL,
17     ↪ task1, NULL);
18 pthread_create(&thread2, NULL,
19     ↪ task2, NULL);
20 ...
21 pthread_join(thread1, NULL);
22 pthread_join(thread2, NULL);
23 }
24 }
```

The code shows a race condition between two threads, `task1` and `task2`, which both manipulate a shared socket `fd`. The `task1` thread performs a `setsockopt` operation, while the `task2` thread performs a `bind` operation. Both threads then call `close(fd)` at the end of their execution. The diagram illustrates that when `task1` calls `close(fd)`, it frees the memory associated with the node, but the node is still part of the linked list because its `next` pointer is still pointing to the next node. This creates a dangling pointer (`dangling ptr`) in the list, which can lead to undefined behavior or security vulnerabilities.

# Challenge 4: No Primitive Needed for Exploitation



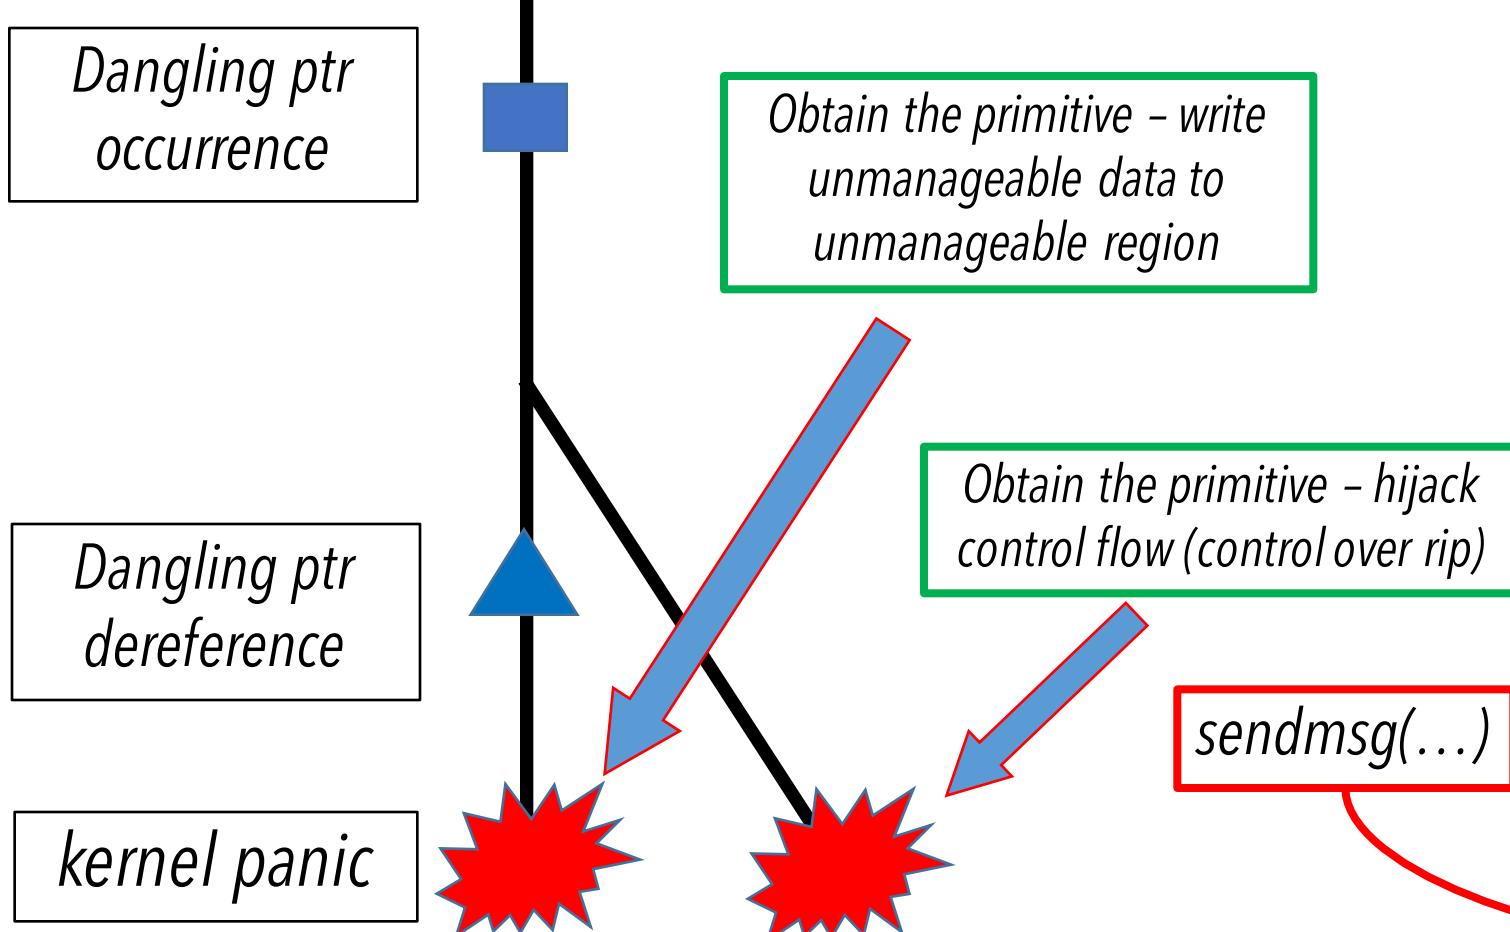
```

1 void *task1(void *unused) {
2 ...
3     int err = setsockopt(fd, 0x107, 18,
4     ↪ ..., ...);
5 }
6 void *task2(void *unused) {
7     int err = bind(fd, &addr, ...);
8 }
9
10 void loop_race() {
11 ...
12     while(1) {
13         fd = socket(AF_PACKET, SOCK_RAW,
14         ↪ htons(ETH_P_ALL));
15 ...
16     //create two racing threads
17     pthread_create(&thread1, NULL,
18     ↪ task1, NULL);
19     pthread_create(&thread2, NULL,
20     ↪ task2, NULL);
21
22     pthread_join(thread1, NULL);
23     pthread_join(thread2, NULL);
24 }

```

A green arrow points from the circled 'setsockopt' call in line 3 to the circled 'task1' call in line 16. Another green arrow points from the circled 'task1' call in line 16 back to the circled 'setsockopt' call in line 3.

# No Useful Primitive == Unexploitable??



```

1 void *task1(void *unused) {
2 ...
3     int err = setsockopt(fd, 0x107, 18,
4     ↪ ..., ...);
5 }
6 void *task2(void *unused) {
7     int err = bind(fd, &addr, ...);
8 }
9
10 void loop_race() {
11 ...
12     while(1) {
13         fd = socket(AF_PACKET, SOCK_RAW,
14         ↪ htons(ETH_P_ALL));
15 ...
16     //create two racing threads
17     pthread_create(&thread1, NULL,
18     ↪ task1, NULL);
19     pthread_create(&thread2, NULL,
20     ↪ task2, NULL);
21
22     pthread_join(thread1, NULL);
23     pthread_join(thread2, NULL);
24 }

```

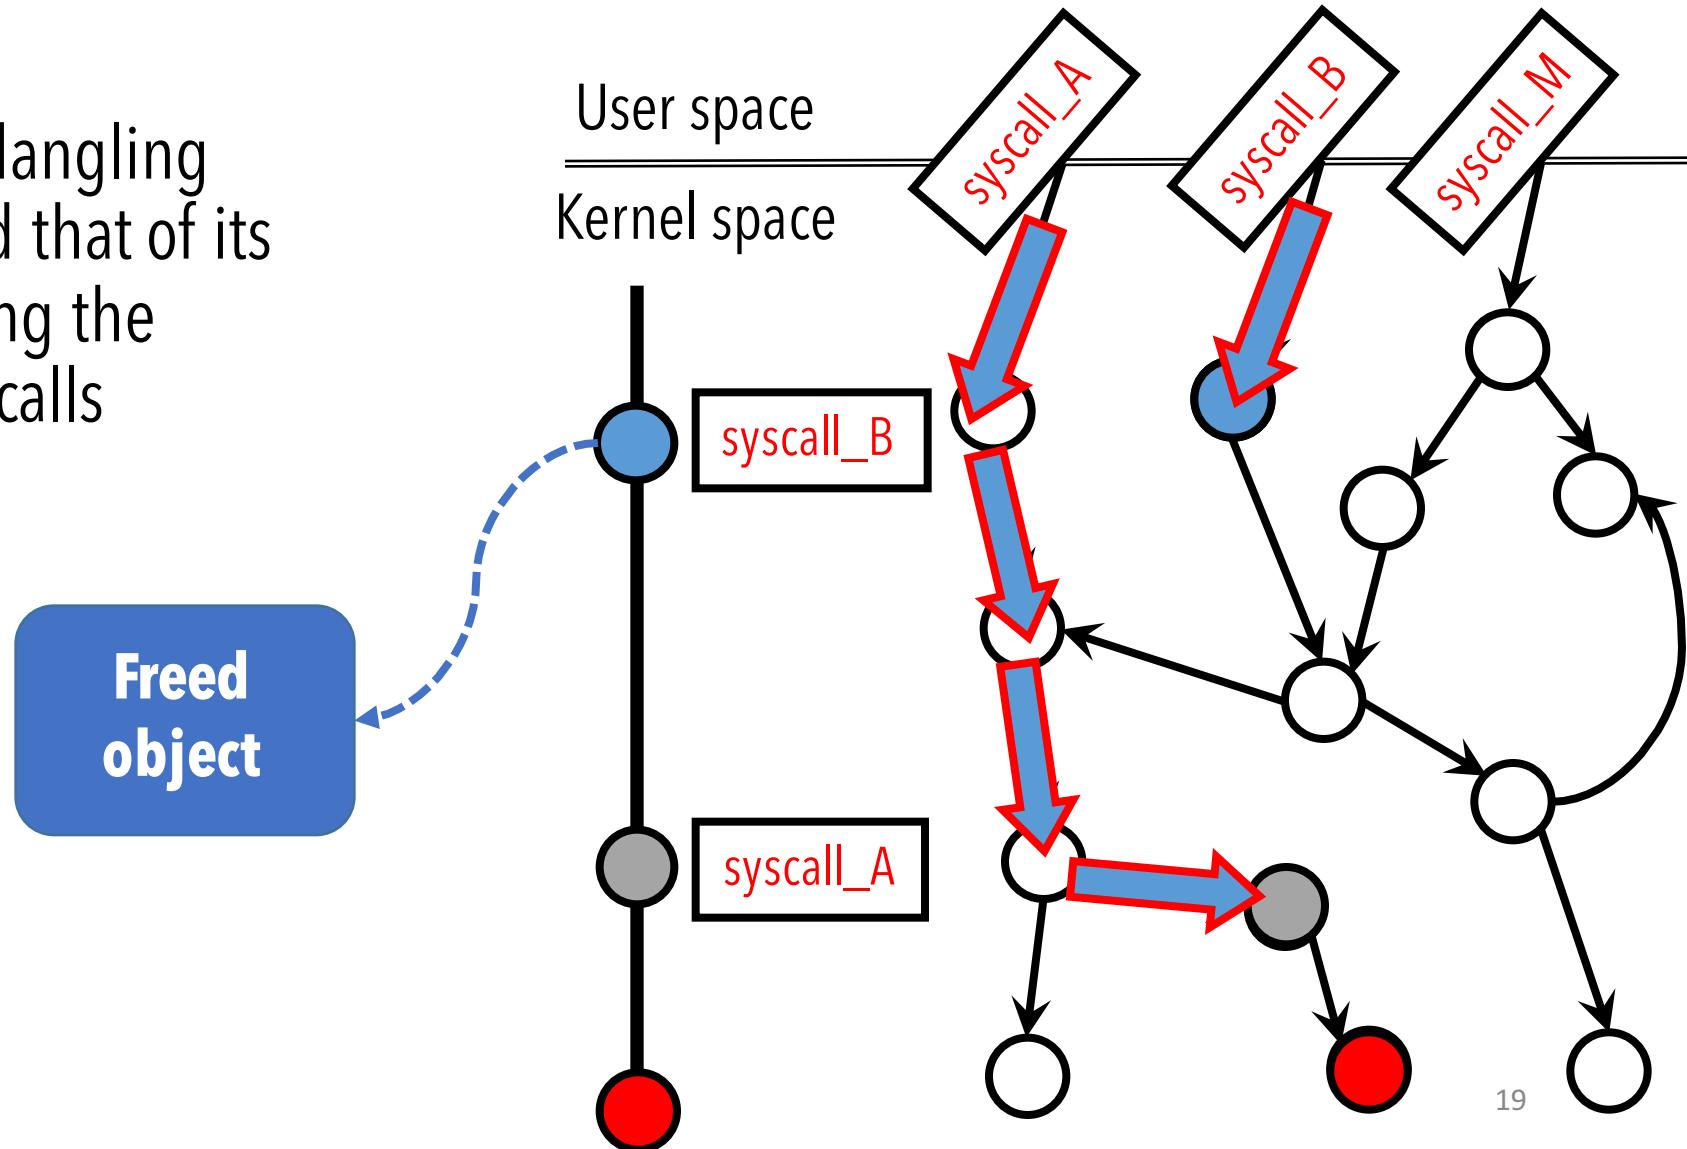
# Roadmap



- Unsolved challenges in exploitation facilitation
- Our techniques - FUZE
- Evaluation with real-world Linux kernel vulnerabilities
- Conclusion

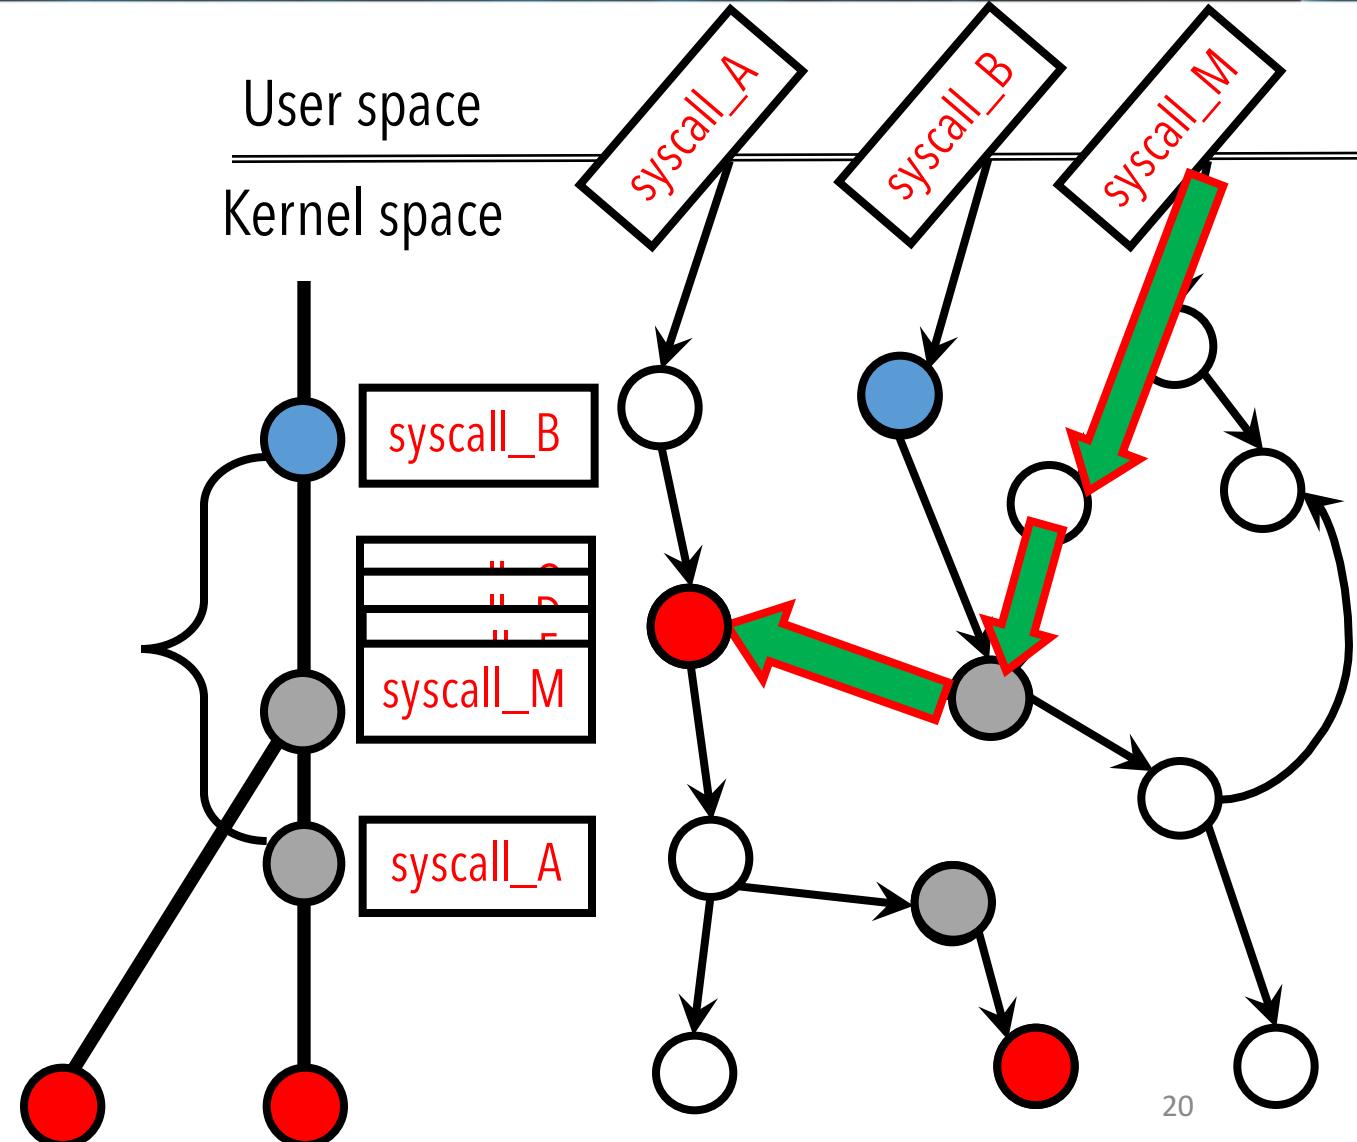
# FUZE - Extracting Critical Info.

- Identifying the site of dangling pointer occurrence, and that of its dereference; pinpointing the corresponding system calls



# FUZE - Performing Kernel Fuzzing

- Identifying the site of dangling pointer occurrence, and that of its dereference; pinpointing the corresponding system calls
- Performing kernel fuzzing between the two sites and exploring other panic contexts (i.e., different sites where the vulnerable object is dereferenced)



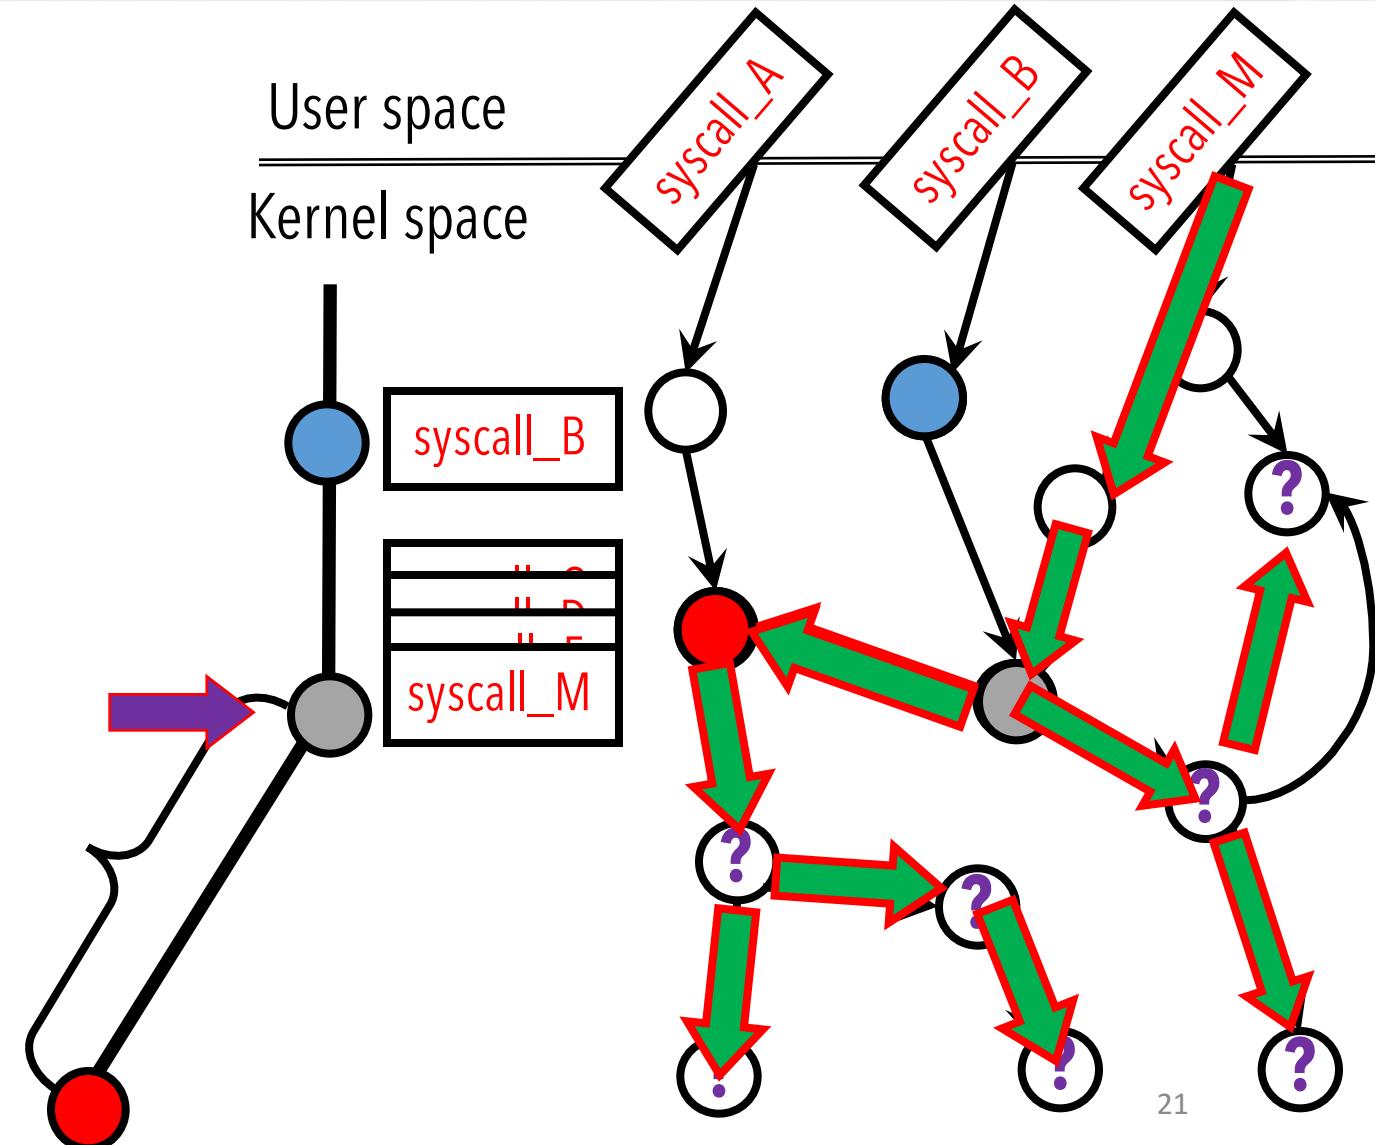
# FUZE – Performing Symbolic Execution

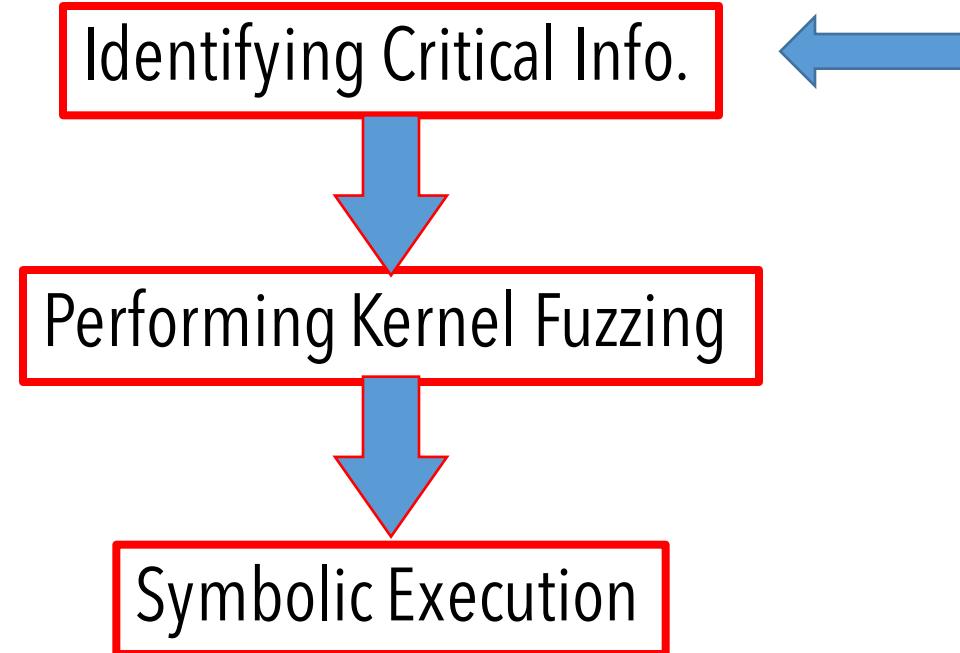
- Identifying the site of dangling pointer occurrence, and that of its dereference; pinpointing the corresponding system calls
- Performing kernel fuzzing between the two sites and exploring other panic contexts (i.e., different sites where the vulnerable object is dereferenced)
- Symbolically execute at the sites of the dangling pointer dereference

**Freed  
object**

Set symbolic value  
for each byte

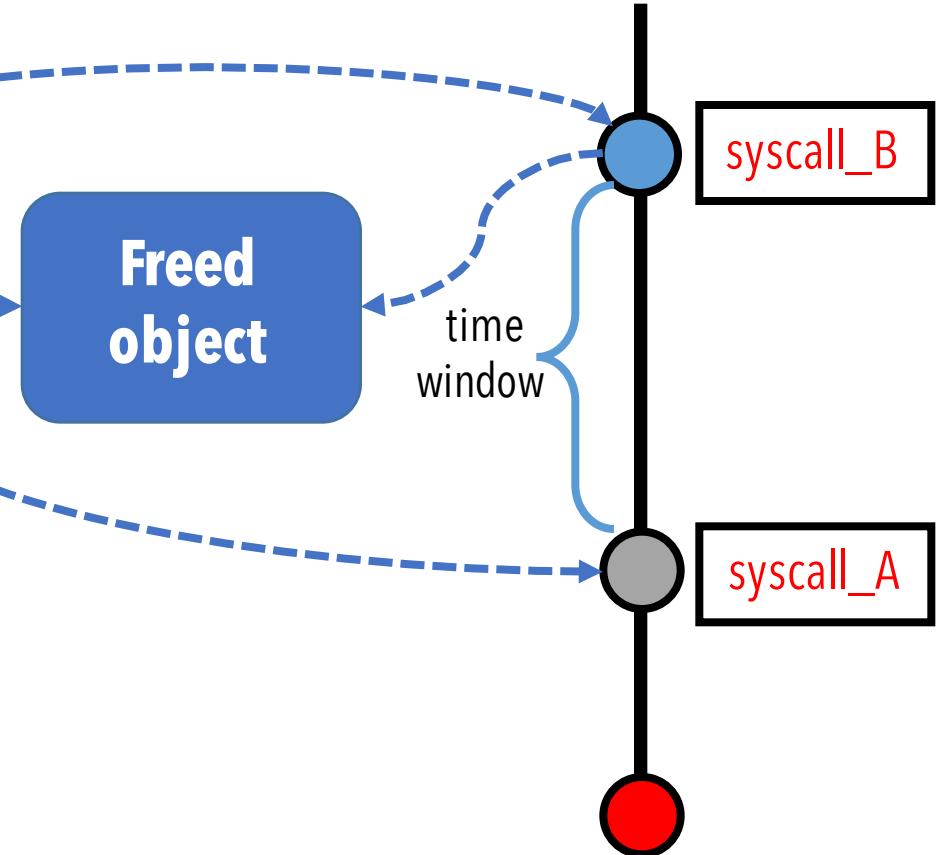
3/9/18





# Critical Information Extraction

- Goal: identifying following critical information
  - Vulnerable object
  - Free site
  - Dereference site
  - Syscalls in PoC tied to corresponding free and dereference
  - Time window between free and dereference
- Methodology:
  - Instrument the PoC with ftrace and generate ftrace log
  - instrument kernel with KASAN
  - Combining both ftrace and KASAN log for analysis



# Critical Information Extraction (cont)

- Goal: identifying following critical information
  - Vulnerable object
  - Free site
  - Dereference site
  - Syscalls in PoC tied to corresponding free and dereference
  - Time window between free and dereference
- Methodology:
  - Instrument the PoC with ftrace[1] and generate ftrace log
  - instrument kernel with KASAN[2]
  - Combining both ftrace and KASAN log for analysis

Unique ID for each syscall in PoC

```
void *task1(void *unused) {  
    ...  
    write_ftrace_marker(1);  
    int err = setsockopt(...);  
    write_ftrace_marker(1);  
}  
void *task2(void *unused) {  
    write_ftrace_marker(2);  
    int err = bind(...);  
    write_ftrace_marker(2);  
}  
...  
void loop_race(){  
    ...  
}  
int main(){  
    ftrace_kmem_trace_enable();  
    loop_race();  
}
```

[1] ftrace. <https://www.kernel.org/doc/Documentation/trace/ftrace.txt>  
8/9/18

[2] kasan. <https://github.com/google/kasan/wiki>

# Critical Information Extraction (cont)

BUG: KASAN: use-after-free  
in **dev\_add\_pack+0x304/0x310**  
Write of size 8 at addr  
**ffff88003280ee70**  
by task poc/**2678**

Call Trace:

...  
Allocated by task **7271**:  
... (allocation trace)  
Freed by task **2678**:  
... (free trace)

The buggy address belongs  
to the object at  
**ffff88003280e600**

which belongs to the cache  
kmalloc-4096 of size 4096

pid:2678

pid:7271

8/9/18  
pid:7272

...  
poc-7271 : tracing\_mark\_write: executing syscall: **setsockopt**  
poc-7272 : tracing\_mark\_write: executing syscall: **bind**  
poc-7271 : **kmalloc**: call\_site=... ptr=**ffff88003280e600**  
bytes\_req=2176 bytes\_alloc=4352 gfp\_flags=GFP\_KERNEL  
...  
poc-7271 : tracing\_mark\_write: finished syscall: **setsockopt**  
...  
poc-7272 : tracing\_mark\_write: finished syscall: **bind**  
...  
poc-2678 : tracing\_mark\_write: executing syscall: **close**  
poc-2678 : **kfree**: call\_site=... ptr=**ffff88003280e600**  
...  
poc-2678 : tracing\_mark\_write: finished syscall: **close**  
poc-2678 : tracing\_mark\_write: executing syscall: **socket**  
...  
end of ftrace

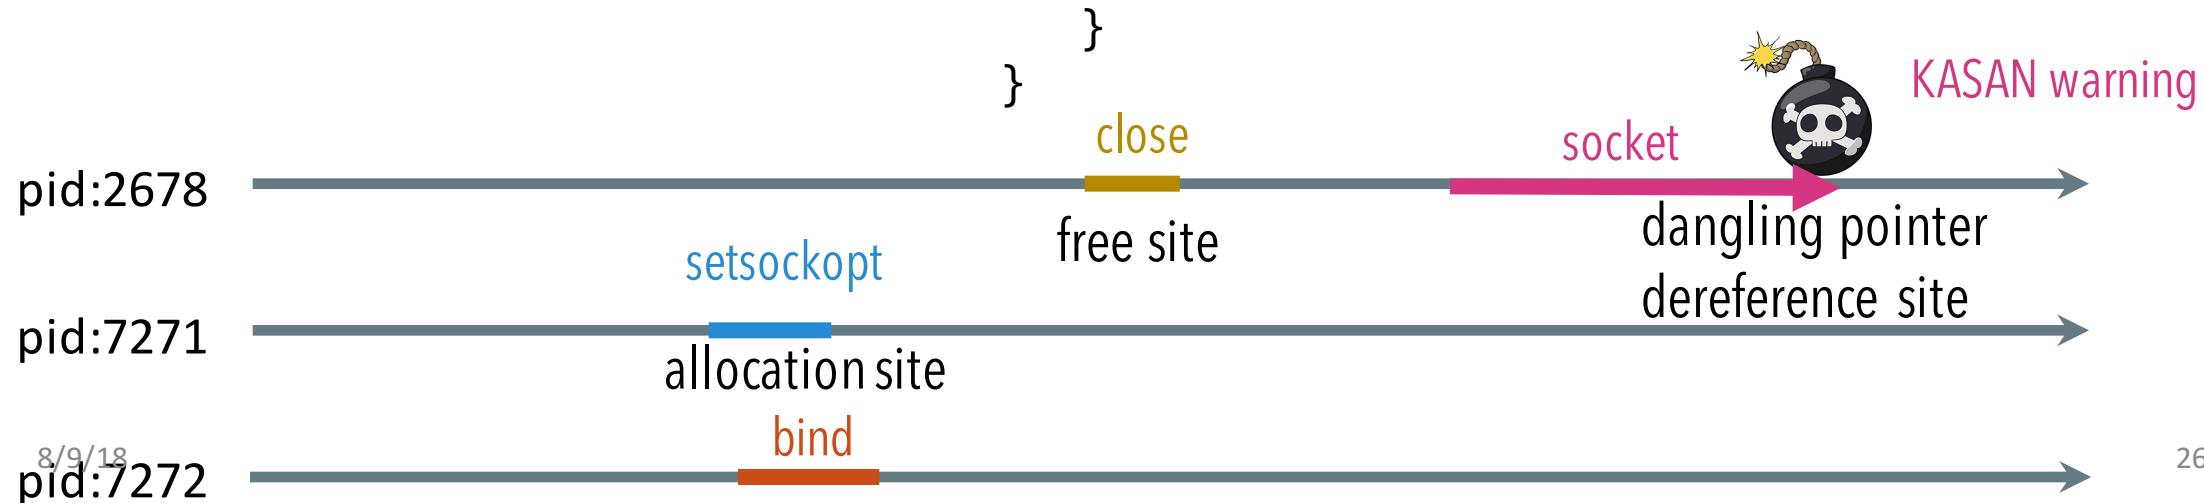


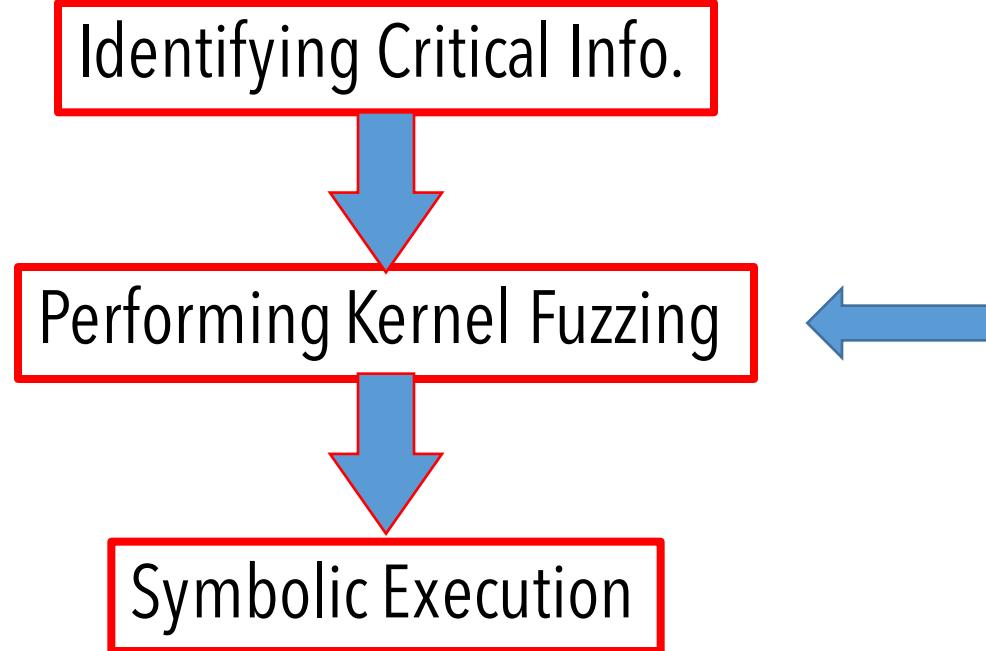
dangling pointer  
dereference site

# Critical Information Extraction (cont)

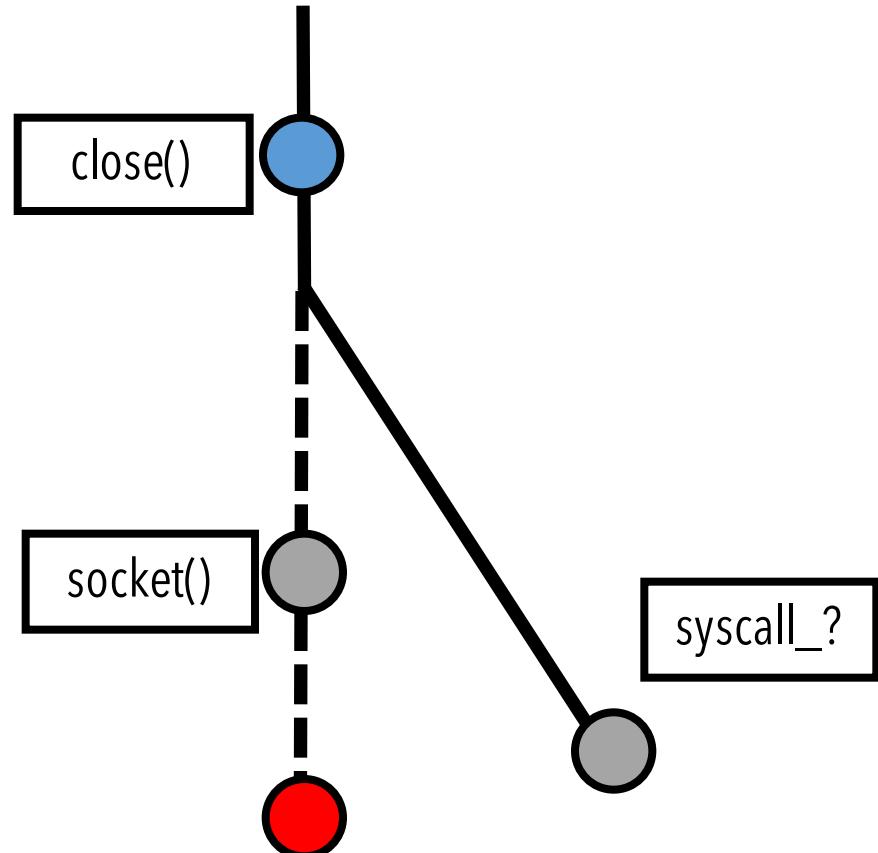
```
void *task1(void *unused) {
    ...
    int err = setsockopt(fd,
0x107, 18, ..., ...);
}
void *task2(void *unused) {
    int err = bind(fd, &addr,
...);
}
```

```
void loop_race() {
    ...
    while(1) {
        fd = socket(AF_PACKET, SOCK_RAW,
htons(ETH_P_ALL));
        ...
        pthread_create (&thread1, NULL, task1, NULL);
        pthread_create (&thread2, NULL, task2, NULL);
        pthread_join(thread1, NULL);
        pthread_join(thread2, NULL);
        close(fd);
    }
}
```





# Kernel Fuzzing

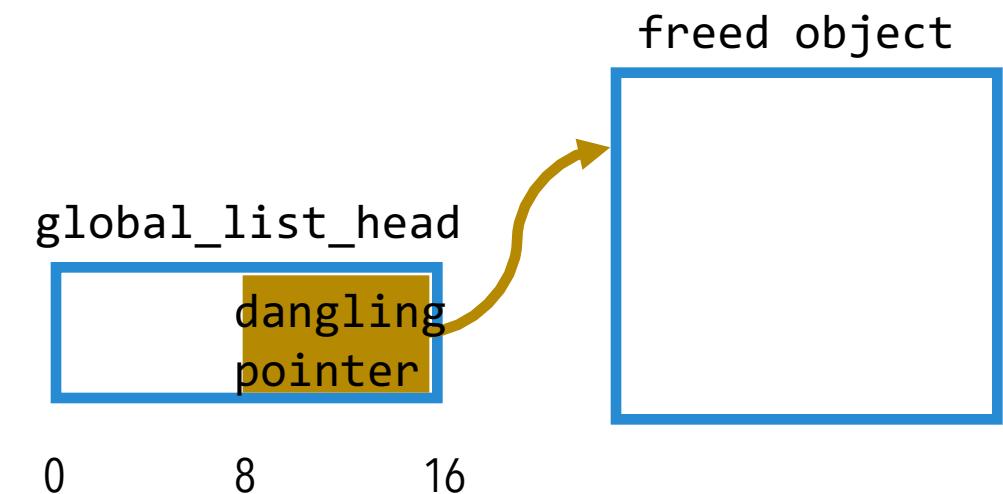


```
poc_wrapper(){  
    /* PoC wrapping function */  
    ...  
    socket(); //dereference site  
    while(true){ // Race condition  
        ...  
        threadA(...);  
        threadB(...);  
        ...  
        close(); //free site  
        /* instrumented statements */  
        if (!ioctl(...)) // interact with  
            a kernel module  
            return;  
    }  
}  
poc_wrapper();  
fuzzing();
```

# Kernel Module for Dangling Pointer Identification

- Identifying dangling pointer through the global variable pertaining to vulnerable object
  - Setting breakpoint at syscall tied to the dangling pointer dereference
  - Executing PoC program and triggering the vulnerability
  - Debugging the kernel step by step and recording dataflow (all registers)
  - Tracking down global variable (or current task\_struct) through backward dataflow analysis
  - Recording the base address the global variable (or current task\_struct) and the offset corresponding to the freed object

```
mov rdx, ds: global_list_head  
...  
mov rax, qword ptr[rdx+8]  
mov rdi, qword ptr[rdx+16] : dang1. deref.
```

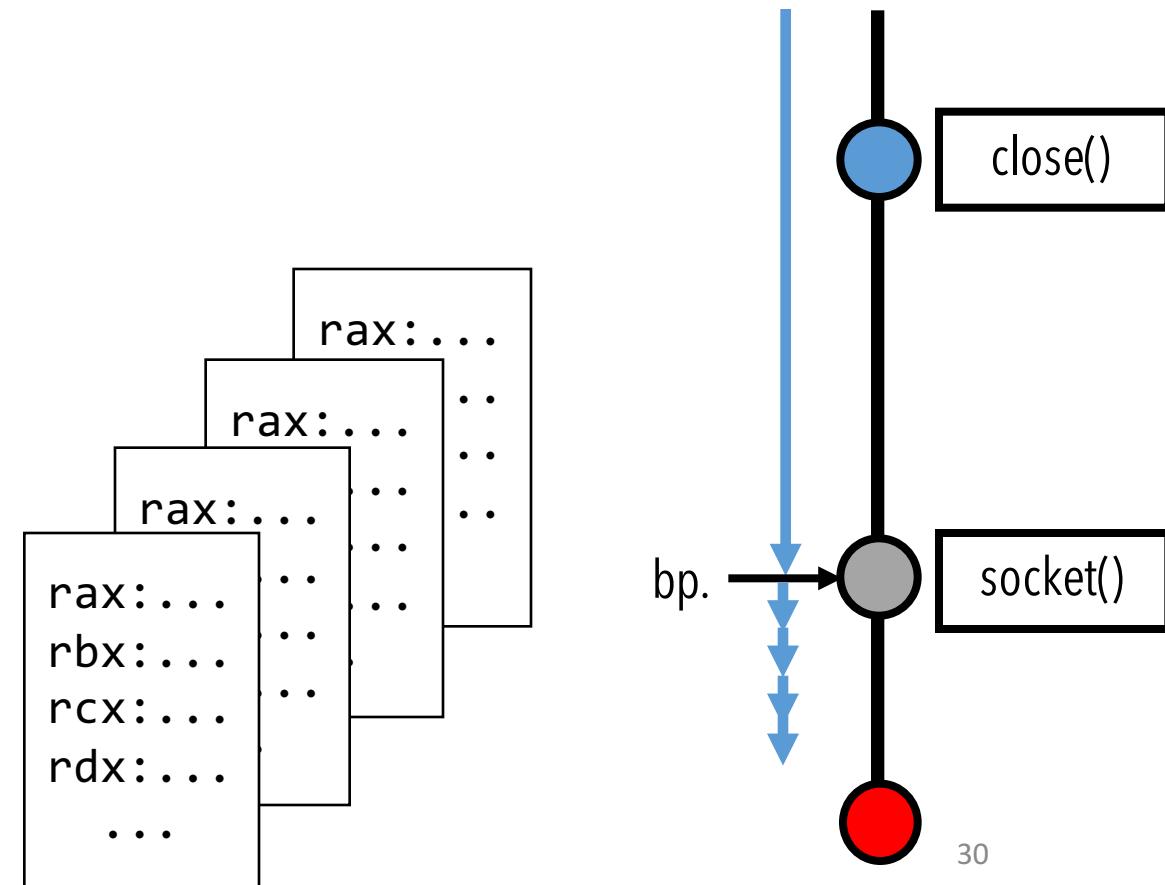


- Identifying dangling pointer through the global variable pertaining to vulnerable object
  - Setting breakpoint at syscall tied to the dangling pointer dereference
  - Executing PoC program and triggering the vulnerability
  - Debugging the kernel step by step and recording dataflow (all registers)
  - Tracking down global variable (or current task\_struct) through backward dataflow analysis
  - Recording the base address the global variable (or current task\_struct) and the offset corresponding to the freed object

```

mov rdx, ds:global_list_head
...
mov rax, qword ptr[rdx+8]
mov rdi, qword ptr[rax+16] : dangl. deref.

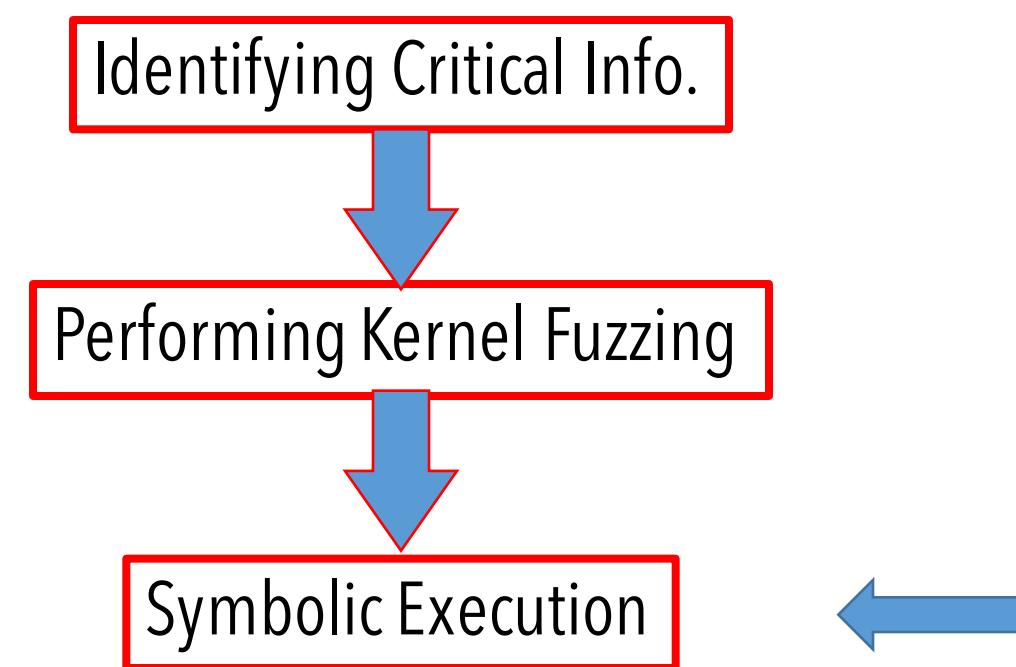
```



# Kernel Fuzzing(cont)

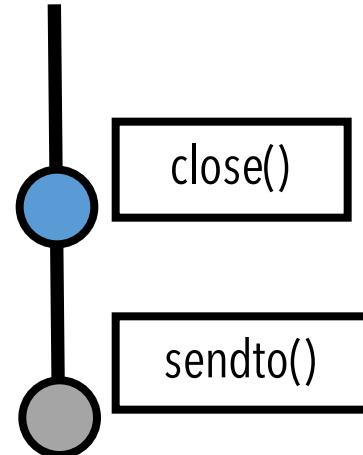
- Reusing syzkaller[1] to performing kernel fuzzing after a dangling pointer is identified
  - generate syz-executor which invoke poc\_wrapper first
- enable syscalls that potentially dereference the vulnerable object
  - "enable\_syscalls"
- transfer variables that appears in the PoC into the interface
  - e.g. file descriptors

```
poc_wrapper();  
fuzzing();
```



# Symbolic Execution

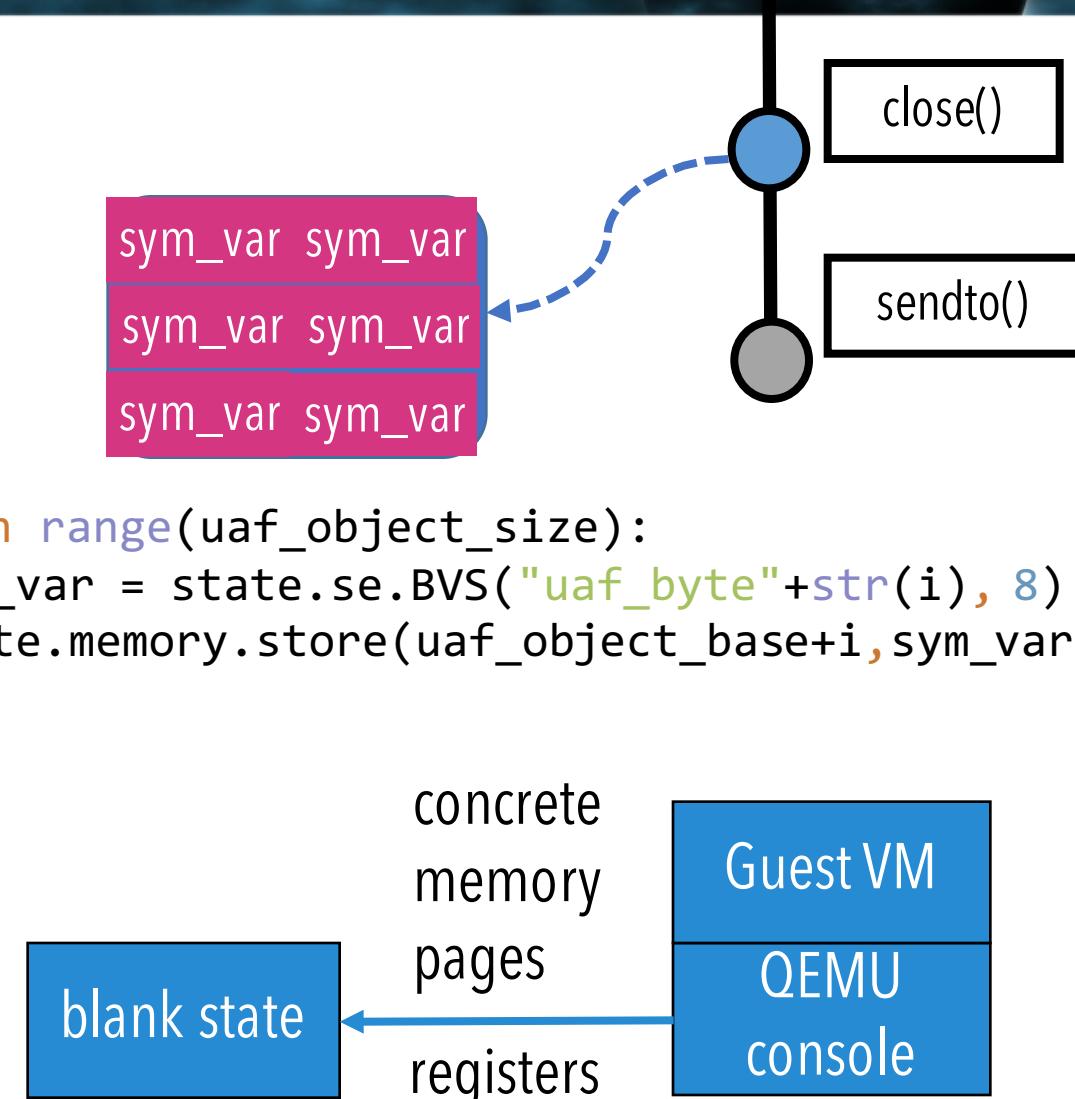
- Symbolic execution for kernel is challenging.
  - How to model and emulate interrupts?
  - How to handle multi-threading?
  - How to emulate hardware device?
- Our goal: use symbolic execution for identifying exploitable primitives
- We can opt-in angr[1] for kernel symbolic execution from a concrete state
  - single thread
  - no interrupt
  - no context switching



# Symbolic Execution

- Symbolic Execution initialization
  - Setting conditional breakpoint at the dangling pointer dereference site
  - Running the PoC program to reach the dangling pointer dereference site
  - Migrating the memory/register state to a blank state
  - Setting freed object memory region as symbolic
  - Starting symbolic execution!

```
for i in range(uaf_object_size):
    sym_var = state.se.BVS("uaf_byte"+str(i), 8)
    state.memory.store(uaf_object_base+i, sym_var)
```

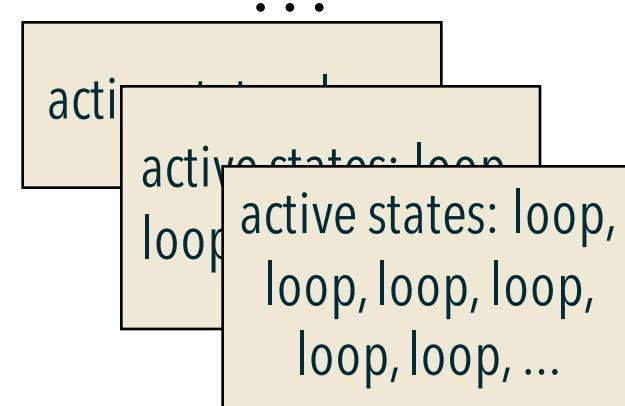


# State(Path) Explosion

Memory consumption  $\approx$  number\_of\_states \* size\_of\_each\_state

- Our design already mitigates state explosion by starting from the first dereference site
  - no syscall issues
  - no user input issues
- However, if a byte from the freed object is used in a branch condition, path explosion occurs.
- Workarounds:
  - limiting the time of entering a loop.
  - limiting the total length of a path.
  - copy concrete memory page on demand
  - kernel function summary

```
    mov edx, dword ptr[freed obj]
loop:
    ...
    inc ecx
    cmp ecx, edx
    jne loop (0xffffffff81abcdef)
```



```
for state in simgr.active:
    if detect_loop(state, 5):
        simgr.remove(state)

for state in simgr.active:
    if len(state.history) > 200:
        simgr.remove(state)
```

# Useful primitive identification

- Unconstrained state
  - state with symbolic Instruction pointer
  - symbolic callback
- double free
  - e.g. `mov rdi, uaf_obj; call kfree`
- memory leak
  - invocation of `copy_to_user` with src point to a freed object
  - syscall return value

Code fragment related to an exploit primitive of CVE-2017-15649

```
if (ptype->id_match)
    return ptype->id_match(ptype, skb->sk)
```

Code fragment related to an exploit primitive of CVE-2017-17053

```
...
kfree(ldt); // ldt is already freed
```

Code fragment related to an exploit primitive of CVE-2017-8824

```
case 127...191:
    return ccid_hc_rx_getsockopt(dp-
>dccps_hc_rx_ccid, sk, optname, len, (u32
__user *)optval, optlen)
```

- write-what-where
  - mov qword ptr [rdi], rsi

rdi (destination)	rsi (source)	primitive
symbolic	symbolic	arbitrary write ( qword shoot)
symbolic	concrete	write fixed value to arbitrary address
free chunk	any	write to freed object
x(concrete)	x(concrete)	self-reference structure
metadata of freed chunk	any	meta-data corruption

- When you found a cute exploitation technique, why not make it reusable?
- Each technique can be implemented as state plugins to angr.
- Exploit technique database
  - Control flow hijack attacks:
    - pivot-to-user
    - turn-off-smap and ret-to-user
    - set\_rw() page permission modification
    - ...
  - Double free attacks
    - auxiliary victim object
    - loops in free pointer linked list
  - memory leak attacks
    - leak sensitive information (e.g. credentials)
  - write-what-where attacks
    - heap metadata corruption
    - function-pointer-hijack
    - vdso-hijack
    - credential modification
    - ...

# From Primitive to Exploitation: SMEP bypass

- Solution: ROP
    - stack pivot to userspace [1]
- control flow hijack  
primitive

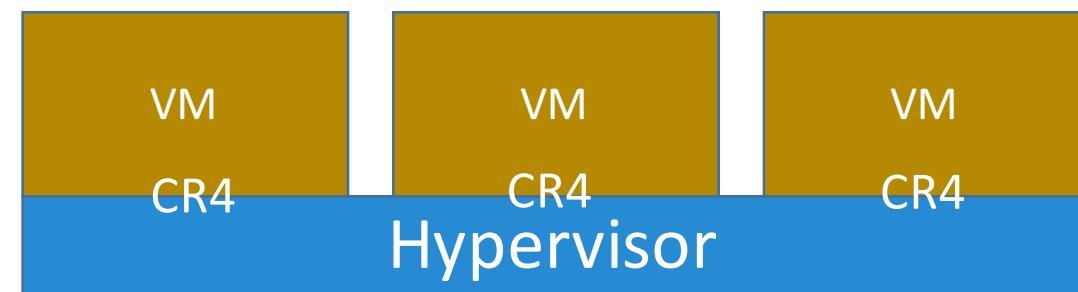
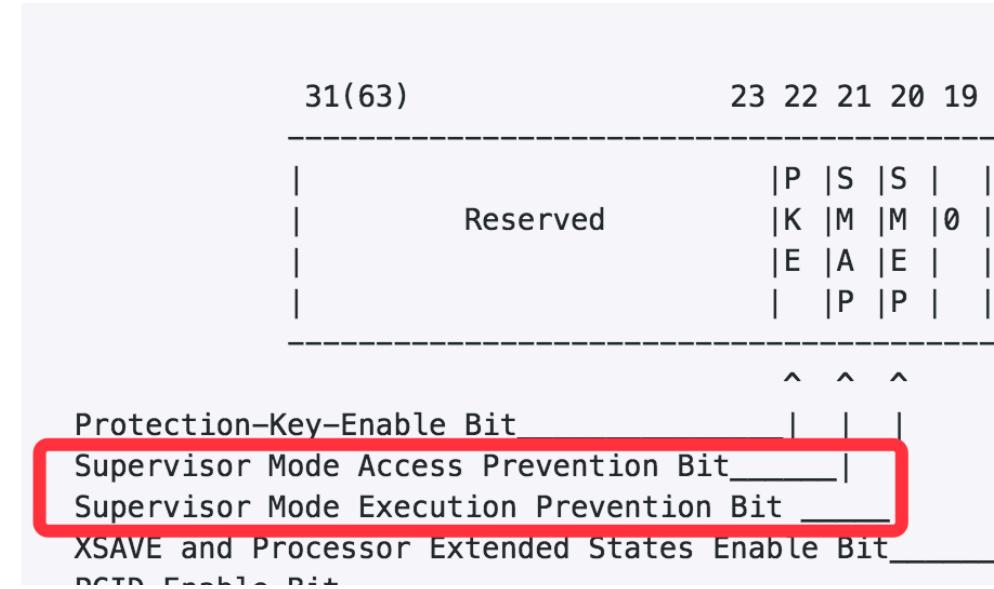
```
mov rax, qword ptr[evil_ptr]  
call rax
```

```
If simgr.unconstrained:  
    for ucstate in simgr.unconstrained:  
        try_pivot_and_rop_chain(ucstate)
```

stack pivot gadget

```
xchg eax, esp ; ret
```

- Solution: using two control flow hijack primitives to clear SMAP bit (21th) in CR4 and land in shellcode
  - 1<sup>st</sup> ... > mov cr4, rdi ; ret
  - 2<sup>nd</sup> ... > shellcode
- limitation
  - can not bypass hypervisor that protects control registers
- Universal Solution: kernel space ROP
  - bypass all mainstream mitigations.



# Extra Symbolization

- Goal: enhance the ability to find useful primitives
- Observation: we can use a ROP/JOP gadget to control an extra register and explore more state space
- Approach:
  - forking states with additional symbolic register upon symbolic states
  - We may explore more states by adding extra symbolic registers

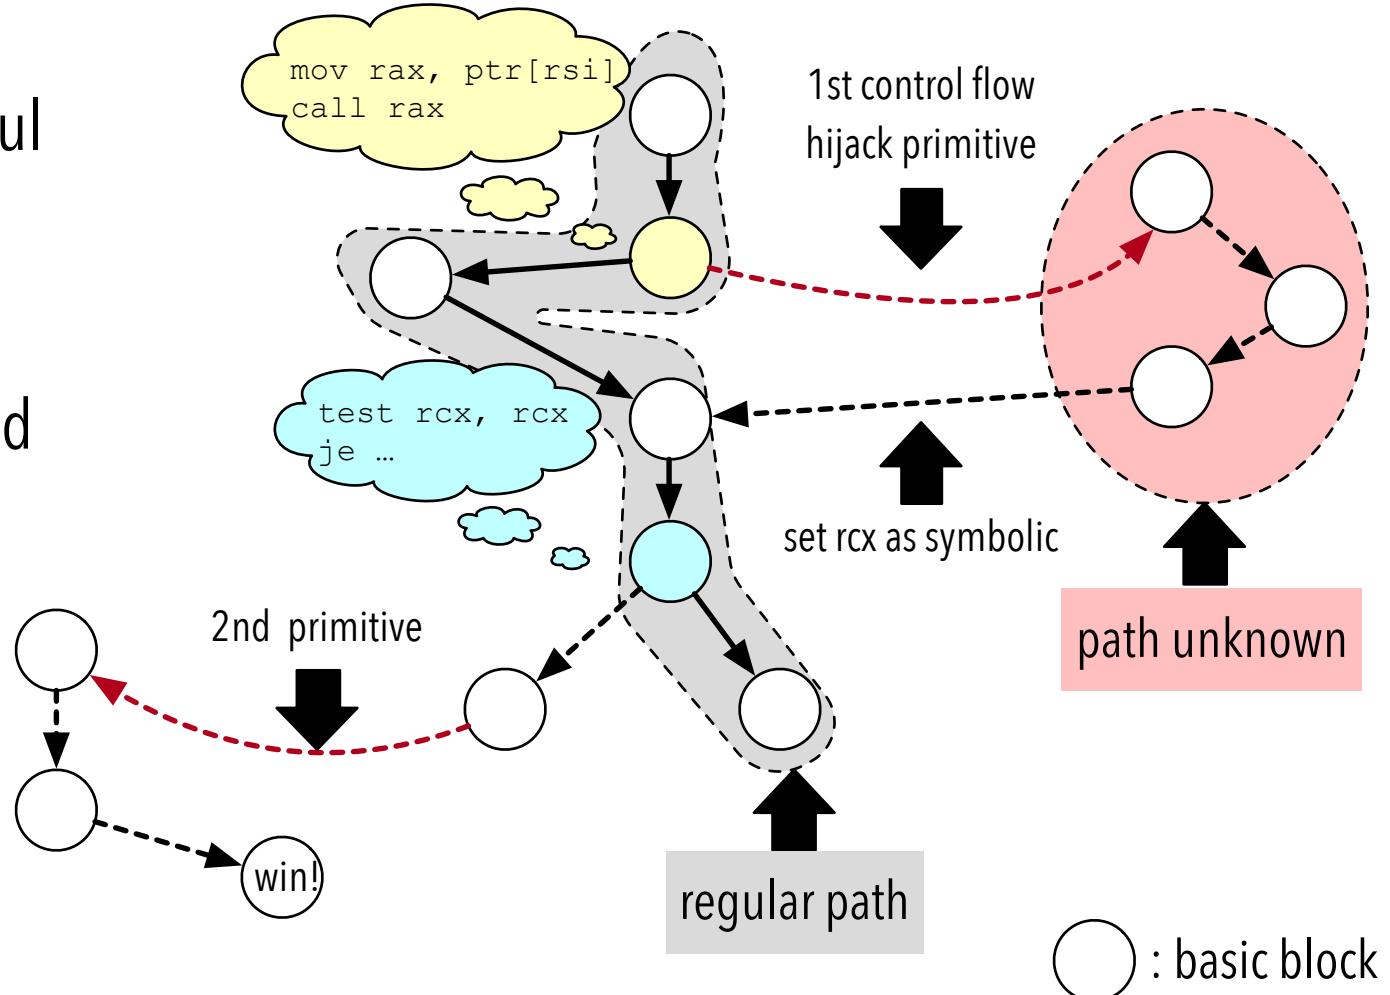


Figure: Identifying two control flow hijack primitive for CVE-2017-15649

# From Primitive to Exploitation: post-exploit fix

- Sometimes we get control flow hijack primitive in interrupt context.
  - avoiding double fault: keep writing to your ROP payload page to keep it mapped in
- Some syscall (e.g. execve) checks current execution context (e.g. via reading preempt\_count) and decides to panic upon unmatched context.

```
BUG_ON(in_interrupt());
```



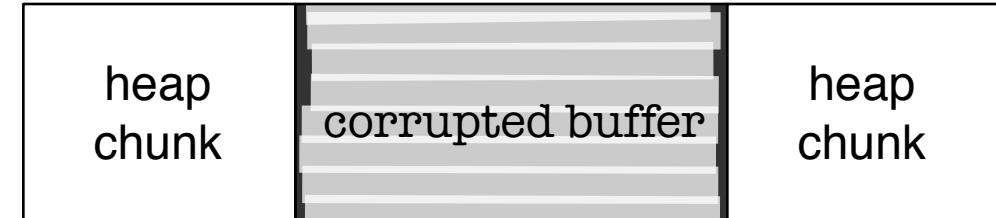
-----[ cut here ]-----  
kernel BUG at linux/mm/vmalloc.c:1394!

- Solution: fixing preempt\_count before invoking execve("/bin/sh", NULL, NULL)

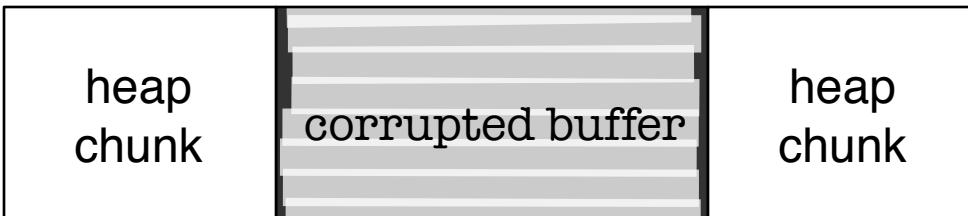
# Symbolic Read/Write

```
t0
    mov rdi, QWORD PTR [corrupted_buffer]
t1
    mov rax, QWORD PTR [rdi]
t2
```

rdi: xxx

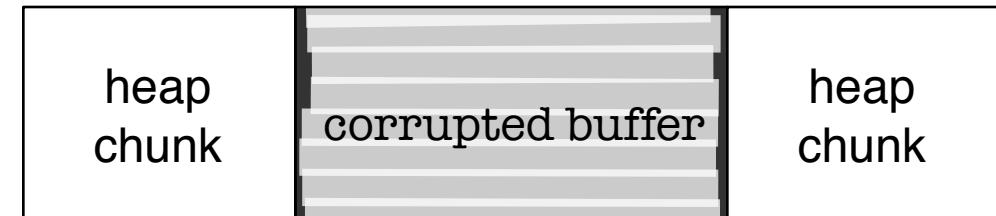


rdi → symbolic\_qword



t1

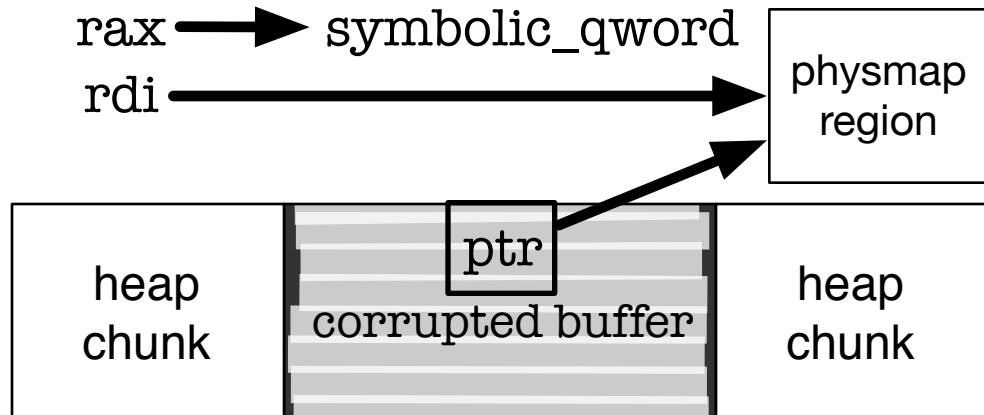
rax → symbolic\_qword  
 rdi → ???



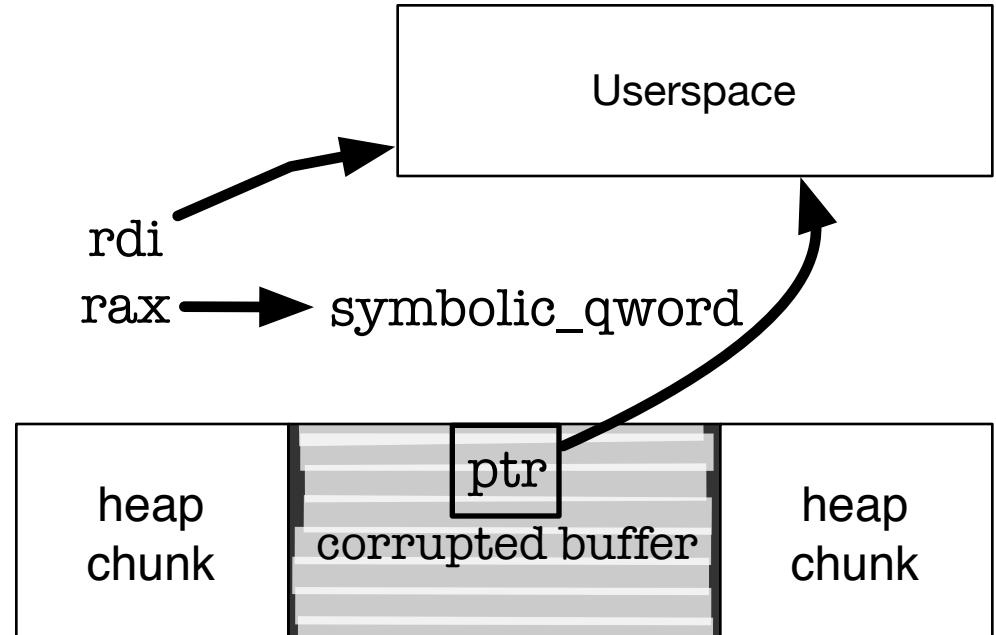
t2

# Symbolic read/write concretization strategy

- Concretize the symbolic address to pointing a region under our control
  - no SMAP: entire userspace
  - with SMAP but no KASLR: physmap region
  - with SMAP and KASLR: ... need a leak first



```
mov rdi, QWORD PTR [corrupted_buffer]
mov rax, QWORD PTR [rdi]
```



# Roadmap



- Unsolved challenges in exploitation facilitation
- Our techniques -- FUZE
- Evaluation with real-world Linux kernel vulnerabilities
- Conclusion

# Case Study

- 15 real-world UAF kernel vulnerabilities
- Only 5 vulnerabilities have demonstrated their exploitability against SMEP
- Only 2 vulnerabilities have demonstrated their exploitability against SMAP

CVE-ID	# of public exploits		# of generated exploits	
	SMEP	SMAP	SMEP	SMAP
2017-17053	0	0	1	0
2017-15649	0	0	3	2
2017-15265	0	0	0	0
2017-10661	0	0	2	0
2017-8890	1	0	1	0
2017-8824	0	0	2	2
2017-7374	0	0	0	0
2016-10150	0	0	1	0
2016-8655	1	1	1	1
2016-7117	0	0	0	0
2016-4557	1	1	4	0
2016-0728	1	0	3	0
2015-3636	0	0	0	0
2014-2851	1	0	1	0
2013-7446	0	0	0	0
overall	5	2	19	46 5

# Case Study (cont)

- FUZE helps track down useful primitives, giving us the power to
  - Demonstrate exploitability against SMEP for 10 vulnerabilities
  - Demonstrate exploitability against SMAP for 2 more vulnerabilities
  - Diversify the approaches to performing kernel exploitation
    - 5 vs 19 (SMEP)
    - 2 vs 5 (SMAP)

CVE-ID	# of public exploits		# of generated exploits	
	SMEP	SMAP	SMEP	SMAP
2017-17053	0	0	1	0
2017-15649	0	0	3	2
2017-15265	0	0	0	0
2017-10661	0	0	2	0
2017-8890	1	0	1	0
2017-8824	0	0	2	2
2017-7374	0	0	0	0
2016-10150	0	0	1	0
2016-8655	1	1	1	1
2016-7117	0	0	0	0
2016-4557	1	1	4	0
2016-0728	1	0	3	0
2015-3636	0	0	0	0
2014-2851	1	0	1	0
2013-7446	0	0	0	0
overall	5	2	19	475

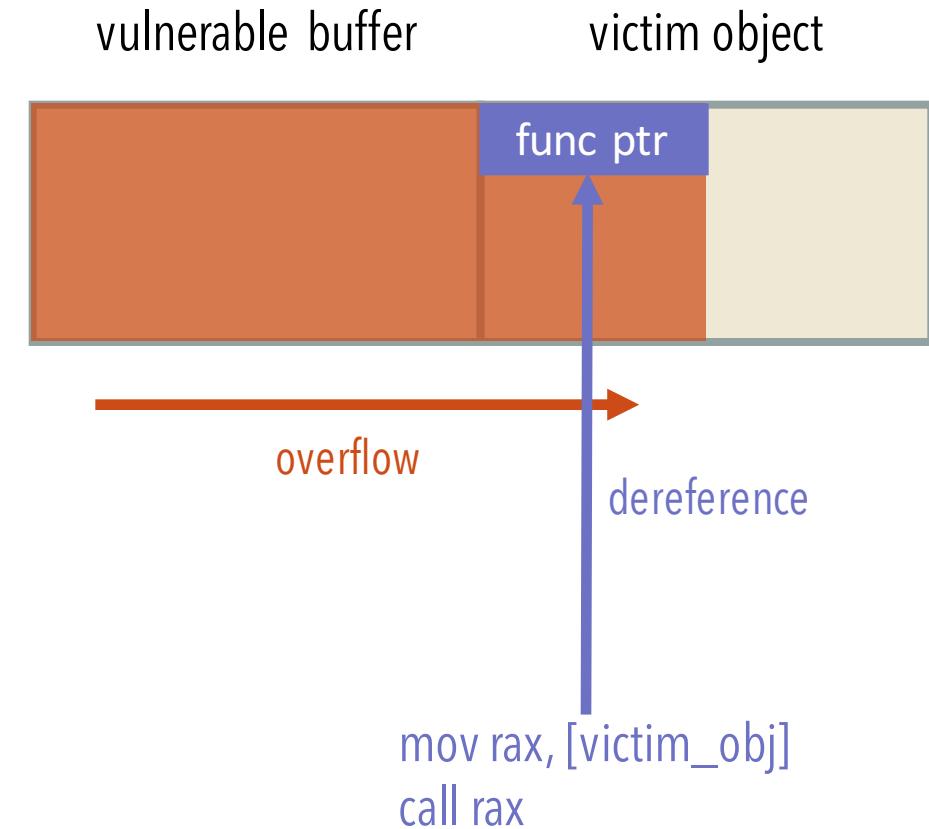
# Discussion on Failure Cases



- Dangling pointer occurrence and its dereference tie to the same system call
- FUZE works for 64-bit OS but some vulnerabilities demonstrate its exploitability only for 32-bit OS
  - E.g., CVE-2015-3636
- Perhaps unexploitable?
  - CVE-2017-7374 ← null pointer dereference
  - E.g., CVE-2013-7446, CVE-2017-15265 and CVE-2016-7117

# What about heap overflow

- Heap overflow is similar to use-after-free:
  - a victim object can be controlled by attacker by:
    - heap spray (use-after-free)
    - overflow (or memory overlap incurred by corrupted heap metadata)
- Heap overflow exploitation in three steps:
  - 1) Understanding the heap overflow  
off-by-one? arbitrary length? content controllable?
  - 2) Find a suitable victim object and place it after the vulnerable buffer  
automated heap layout[1]
  - 3) Dereference the victim object for exploit primitives



[1] Heelan et al. Automatic Heap Layout Manipulation for Exploitation. USENIX Security 2018.



**236.5 million**

Largest retailer in China,  
online or offline shoppers



**\$37.5bn**

Third largest internet company  
in the world by revenue in 2016



First e-commerce company to use commercial drone delivery

## 700 Million

June Sales Event **Items Sold**

Massive Scale

## 236.5M

active customer accounts

## 120K

active third-party vendors on  
JD platform

## 120K

full-time employees

## 1.59B

full-time orders fulfilled in 2016



- Bug prioritization
  - Focus limited resources to fix bugs with working exploits
- APT detection
  - Use generated exploits to generate fingerprints for APT detection
- Exploit generation for Red Team
  - Supply Red Team with a supply of new exploits

# Roadmap



- Unsolved challenges in exploitation facilitation
- Our techniques -- FUZE
- Evaluation with real-world Linux kernel vulnerabilities
- Conclusion

# Conclusion



- Primitive identification and security mitigation circumvention can greatly influence exploitability
- Existing exploitation research fails to provide facilitation to tackle these two challenges
- Fuzzing + symbolic execution has a great potential toward tackling these challenges
- Research on exploit automation is just the beginning of the GAME! Still many more challenges waiting for us to tackle...

- Acknowledgement:
  - Yueqi Chen
  - Jun Xu
  - Xiaorui Gong
  - Wei Zou
- Exploits available at:
  - [https://github.com/ww9210/Linux\\_kernel\\_exploits](https://github.com/ww9210/Linux_kernel_exploits)
- Contact: [wuwei@iie.ac.cn](mailto:wuwei@iie.ac.cn)



- Acknowledgement:
  - Yueqi Chen
  - Jun Xu
  - Xiaorui Gong
  - Wei Zou
- Exploits available at:
  - [https://github.com/ww9210/Linux\\_kernel\\_exploits](https://github.com/ww9210/Linux_kernel_exploits)